

JPL-SIP Report

Introducing Covariance Constraints to Atmospheric Control Adjustments Generated Via Ocean Climate Simulation

Shauntclair Ruiz

Mentored by Ian Fenty, Ou Wang at Jet Propulsion Laboratory

August 19th, 2024

Contents

1 Abstract	5
2 Background	5
2.1 MITgcm: Ocean Climate Simulation	5
2.2 MITgcm: Atmospheric Adjustments	6
2.3 MITgcm: Gradient and the Adjoint Method	6
2.4 Gradient Descent	7
3 Conventions	8
3.1 Important Variables	8
3.2 Iteration	8
3.3 Gradient Descent	9
3.4 Loss Function	9
3.5 Gradient dJ/df	9
4 Introduction to our Work	11
4.1 Our Problem	11
4.2 Our Solution: Covariance Constraints	11
4.2.1 Benefits of Covariance Constraints	11
4.3 Measuring Covariance Similarity	12
4.4 Various Covariance Constraint Methods	12
4.4.1 Modifying the loss function J	12
4.4.2 Dan's Method	13
4.4.3 Dan's Method Modified	13
5 Methods	15
5.1 Discretizing our state	15
5.2 Advection-Diffusion-Forcing Model	15
5.2.1 Continuous Differential Equation	15
5.2.2 Discretization of Differential Equation	15
5.3 World Parameter Generation	16
5.3.1 Initial State	16
5.3.2 Atmosphere	16
5.3.3 Current Flow	17
5.3.4 Other parameters	18
5.3.5 Cyclic Variables	18
5.4 Observations	18
5.4.1 $Hy(t)$ notation	18
5.4.2 Unobserved grid cells	18
5.5 Computing the gradient	19
5.6 Implementing our gradient descent variants	19
5.7 Measuring performance: success metrics	19
5.7.1 Ocean Misfit	19

5.7.2	Atmoshere Misfit	19
5.7.3	Mahalanobis Distance	20
6	Results	21
6.1	Simple Gradient Descent vs. Modifying J	21
6.1.1	Plots	21
6.1.2	Conclusions	23
6.2	Dan Method	23
6.2.1	Plots	24
6.2.2	Conclusions	26
6.3	Dan Method Modified	26
6.3.1	Plots	26
6.3.2	Conclusions	27
6.4	Why does the "Dan Modified" Method Fail?	28
6.4.1	PCA Analysis	28
6.4.2	"Forgetting" theory	30
6.4.3	Supporting evidence: $-Cs_i \cdot \alpha_i$	31
6.4.4	PCA for comparing methods	33
6.4.5	Switching Methods	34
6.5	Is the Dan method really flawed at all?	37
6.5.1	Mahalanobis pressures Correlation, not Covariance	38
6.5.2	The Mahalanobis Distance we expect	39
6.5.3	Conclusion	40
7	Future Work	41
7.1	Extensions of our Toy Model	41
7.2	Extensions Beyond our Toy Model	41
7.3	Further investigation on our work	41
8	References	43
9	Acknowledgments	44
10	Appendices	45
10.1	Matrix Conventions	45
10.2	Dan's Method Derivation	46
10.2.1	Constraints of Dan's Method	46
10.2.2	Optimizing over our constraints	47
10.2.3	Derivation	47
10.3	Dan's Modified Method Derivation	49
10.3.1	Optimizing over constraints	49
10.3.2	Derivation	50
10.4	The Advection-Diffusion-Forcing Model, Discretized	51
10.4.1	Diffusion	51
10.4.2	Forcing	51

10.4.3	Advection	51
10.4.4	Affine Model	52
10.5	Our "gaussian decay" covariance assumption	53
10.6	The χ -squared distribution	54
10.6.1	Informal Justification	54
10.6.2	Formally demonstrating this result	54
10.7	The Adjoint Method	56
10.7.1	Setting Up/Motivation	56
10.7.2	Our Model	56
10.7.3	Redundant Calculations	57
10.7.4	The Adjoint Method: Base Case	57
10.7.5	The Adjoint Method: Recursion	58
10.7.6	Using the adjoint	58
10.7.7	Why is the adjoint useful?	59
10.8	Code	61

1 Abstract

Climate scientists need to make the best estimates of the atmosphere possible. One way to evaluate an atmospheric estimate, is to use that estimate in a physics simulation of the ocean, and compute the accuracy of the resulting ocean simulation. Based on how the ocean simulation differs from ocean observations, we can use gradients computed from simulation equations to determine the best way to adjust the atmospheric estimate, to improve ocean simulation accuracy. This method has a weakness, however: just because an atmosphere creates a more accurate ocean simulation, doesn't mean it's a more accurate atmosphere. Some inaccurate atmospheres may create a more realistic ocean simulation. Our solution is to modify our adjustment, and pressure it to have realistic statistical properties. In particular, we constrain the adjustment covariance to be more similar to the realistic, expected covariance. This eliminates some adjustments that, despite improving simulation accuracy, are statistically unrealistic. We test this method by creating an ocean simulation, modeling observations of that simulation, and using gradient descent to optimize the estimated atmosphere based on those observations, comparing performance with and without covariance constraints. Our approach measurably improves accuracy of atmospheric estimates. This technique could be implemented in large-scale climate models.

2 Background

2.1 MITgcm: Ocean Climate Simulation

As climate change worsens, and becomes more difficult to predict, it's ever more important to have effective climate models. This is exactly what our project pursues: an improvement to our existing climate modeling technology.

Here, we iterate on a powerful tool for simulating ocean climate: MITgcm (MIT General Circulation Model). This model divides the ocean state into discrete grid cells on a 2D map, and simulates that ocean state over discrete timesteps, starting from some initial state. This simulation runs according to approximated physical laws, and physical parameters. The output is a prediction for how the ocean state evolves in time, starting from those initial conditions.

- The parameters in the MITgcm model can represent any facet of the ocean that could affect our state: heat diffusion rate between grid cells, ocean current flows, atmospheric conditions, etc.
- For the purposes of this paper, we'll focus on one subset of parameters: atmospheric conditions.

This, however, is only the "forward-time" process of MITgcm. Next, we consider the "reverse-time" process of MITgcm.

2.2 MITgcm: Atmospheric Adjustments

We've already discussed one application of this model: predicting possible future ocean climates, given different possible atmospheric conditions. However, we can also use MITgcm to improve our estimates of real-world atmospheric conditions. How?

We'll assume our simulation physics are accurate. If so, an accurate atmosphere should generate an accurate ocean simulation.

However, the atmosphere we provide to MITgcm is only an estimation, based on observations. This estimate is limited by the quality and quantity of those observations. Uncertainty in those observations creates an imperfect atmospheric estimate.

MITgcm receives this atmospheric estimate, and simulates the ocean over time, starting from some initial state.

Inaccuracies in our atmospheric estimates cause inaccuracies in our ocean simulation. We can measure this by comparing that ocean simulation to ocean observations.

This is actually something we can use to our advantage: we use our observations as a proxy for the "true" state of the ocean, and we can measure how far off our simulation is.

- We would expect a more accurate atmospheric estimate to create a more accurate ocean state prediction. So we try the converse: if we find an atmospheric estimate that better predicts real-life ocean state, it might be closer to real-life atmospheric conditions.

We assume, however, that our atmospheric estimate is already close to correct. So, rather than creating a new atmosphere from scratch, we add an **adjustment** to our atmospheric estimate.

2.3 MITgcm: Gradient and the Adjoint Method

How do we select our adjustment? We want to adjust our atmosphere, in order to improve the accuracy of our ocean simulation.

- Our atmosphere will be represented as a vector f .
- The "inaccuracy" of our ocean simulation will be represented by the loss function J .

Based on these conventions, we want to determine the way of modifying f that will maximally decrease J (decrease inaccuracy). The most useful piece of information for adjusting f is the gradient, dJ/df .

- This gradient tells us the direction that *locally* increases J by the maximal amount.

Once we have this gradient, we can use it to modify atmosphere f , and improve our ocean simulation accuracy.

- This adjusted atmosphere f_{new} more accurately predicts ocean conditions: thus, we expect it to be more accurate.

In the most simple approach, computing dJ/df would involve a very large multivariable chain rule. However, in a massive model like MITgcm, this is prohibitively expensive.

Instead, MITgcm uses a more efficient method for computing gradients, known as the **adjoint method**. The details are outlined in [10.7](#)

2.4 Gradient Descent

In this paper, we'll put aside the more sophisticated optimization methods that MITgcm uses. Instead, we'll use gradient descent to optimize our atmosphere f .

This means we iteratively apply repeated updates to our atmospheric estimate. After each estimate, we re-compute dJ/df , and take another update.

In the simplest case, using a step size $\eta \in \mathbb{R}$, our update rule takes the form:

$$f_{\text{new}} = f_{\text{old}} - \eta \left(\frac{dJ}{df} \right)_{f=f_{\text{old}}} \quad (1)$$

3 Conventions

3.1 Important Variables

Here, we introduce some conventions useful for the rest of this paper:

The **control vector** f , also known as the **atmospheric control**.

- We're using our the atmosphere to determine, or "control", the simulated ocean state.
- Thus, we call this our **control** vector.

The **control adjustment** a .

- Our optimization will aim to improve f , "adjusting" it with some other vector.
- We call this vector our **control adjustment**:

The **update** u .

- We'll actually be adjusting f multiple times, using gradient descent.
- Each of these is an **update** to our control.

3.2 Iteration

This notation is still insufficient: we need a distinct update u for every iteration of gradient descent.

We'll address this by using i to indicate the particular iteration of gradient descent we are discussing:

- u_i is the gradient descent update that occurs during iteration i .

This convention can be extended to our other variables, f and a :

- f_i is our current control estimate at the beginning of iteration i , before we add u_i .

$$f_{i+1} = f_i + u_i \quad (2)$$

- a_i is our total control adjustment at the beginning of iteration i : it's the sum of every gradient descent step we've taken so far.

$$a_i = \sum_{j=0}^{i-1} u_j \quad (3)$$

We can represent our current control as the sum of our "first-guess" control f_0 , plus our total control adjustment, a_i .

$$f_i = f_0 + a_i \quad (4)$$

3.3 Gradient Descent

Below, we'll consider several "modified" version of gradient descent. All of them involve modifying the update rule, u_i .

But for the simplest version of gradient descent, we can define our update rule:

$$u_i = -\eta \left(\frac{dJ}{df} \right)_{f=f_i} \quad (5)$$

3.4 Loss Function

Above, we mentioned a loss function J , being our measure of "inaccuracy" in our ocean simulated.

How do we measure this? We'll take a common approach: we measure the **squared difference** between the ocean state we want to match, and the ocean state we simulated. We'll repeat this at each timestep.

- The "ocean state we want to match" is the **observed** ocean. We'll represent our observations at time t as $H_y(t)$ (the justification for this notation is provided in [5.4.1](#)).
- The ocean state we simulate for time t will be represented as $x(t)$.

$x(t)$ and $H_y(t)$ are vectors of the same length. So, in order to get the squared difference, we'll need to multiply them as:

$$J_t(x) = \left(x(t) - H_y(t) \right)^T \left(x(t) - H_y(t) \right) \quad (6)$$

Finally, we add up these "misfits" J_t over all timesteps. Our final timestep will be designated as τ .

$$J(x) = \sum_{t=0}^{\tau} \left(x(t) - H_y(t) \right)^T \left(x(t) - H_y(t) \right) \quad (7)$$

3.5 Gradient dJ/df

This loss function J is what we wish to optimize over.

Our next important property is the gradient dJ/df : this will be a crucial property in every one of our gradient descent methods defined below. We'll need to condense our notation for later equations.

Each timestep will have its own gradient, based on the current control estimate, f_i . This gradient is also known as a "sensitivity". Thus, we notate it as:

$$s_i = \left(\frac{dJ}{df} \right)_{f=f_i} \quad (8)$$

Thus, our gradient descent update rule can be simplified as:

$$u_i = -\eta s_i \quad (9)$$

4 Introduction to our Work

4.1 Our Problem

This technique is useful, but there's a major concern: we would expect the real-life atmosphere to better predict the ocean state, but the converse isn't always true. Just because a chosen atmosphere would induce a realistic ocean state, doesn't mean it's a realistic atmosphere.

- This is especially a problem in the case of MITgcm: we have an enormous number of parameters (roughly 100,000), and not nearly enough observations to fully constraint that model.
- There may be many ways to fit our atmosphere to the ocean observations.

We need more information we can use to constrain the problem.

4.2 Our Solution: Covariance Constraints

Our approach is to use the **statistical structure** of the atmosphere to constrain the problem. In other words, we already know some information about what we expect the atmosphere looks like: we should encourage it to match that pattern.

What kind of structure are we referring to? We'll start by considering an example. We expect nearby regions of the atmosphere to "vary together": if we learn that one region is hot, we expect very nearby regions to be hot, too.

We can represent this with **covariance**: we can use simple assumptions, or past data, to estimate the expected covariance between different regions of the atmosphere.

Based on this, we can state our solution more clearly: when we're optimizing f , we constrain our adjustment to have a covariance close to the expected covariance.

- We'll designate this with a covariance matrix C .

4.2.1 Benefits of Covariance Constraints

We hope to get two major benefits from this approach:

- By constraining our optimization to be more realistic, the resulting solution should be more realistic.
- We can *spatially propagate information*: if we learn information about one cell, and another cell is closely correlated, then we can apply this information to that other cell, too.
 - As a result, we can learn more from each observation.

4.3 Measuring Covariance Similarity

We need one more tool: we want our improvement to have a covariance similar to C , but this requires a *metric* for determining how similar our covariance is to C .

We implement this using the **mahalanobis distance**. For a vector z , our mahalanobis distance is

$$z^T C^{-1} z \quad (10)$$

This function will become larger if z has a covariance more different from C .

However, this function has a secondary effect: it penalizes the magnitude of z . Even if z has the correct correlation structure, increasing the magnitude will increase the mahalanobis distance.

- This effect may act as regularization, so for the time being, we leave it be.

Where does our covariance matrix C come from, and what does it look like? This problem is addressed in [10.5](#).

4.4 Various Covariance Constraint Methods

Here, we'll list out each of the methods we attempted, to enforce covariance constraints.

4.4.1 Modifying the loss function J

In optimization, typically, the simplest way to introduce a constraint is to include it in the loss function.

- So, we'll add the mahalanobis distance to our loss function, creating a new function, J' .
- We want to penalize the covariance dissimilarity of our overall adjustment, a_i .
- We'll include a scaling factor α to determine how strongly we want to prioritize our covariance constraint.

$$J' = J + \alpha \cdot a_i^T C^{-1} a_i \quad (11)$$

With this new loss function, we can take the gradient, and run gradient descent.

$$u_i = -\eta \left(\frac{dJ'}{df} \right)_{f=f_i, a=a_i} \quad (12)$$

$$u_i = -\eta \left(s_i + \alpha \cdot 2C^{-1} a_i \right) \quad (13)$$

4.4.2 Dan's Method

Dan Amrhein [1] proposed a different approach for the optimization: thus, we informally refer to this as "Dan's Method".

In this method, we first compute s_i . Then, we want to choose our update u_i , according to two requirements:

- We want our update rule to improve our ocean misfit J . We'll select some desired decrease in J , δ .

We can approximate the change in J as $s_i^\top u_i$. So, we'll use the constraint:

$$s_i^\top u_i = \delta \quad (14)$$

- We want our update u_i to have a similar covariance to C . So, we want to penalize the mahalanobis distance,

$$u_i^\top C^{-1} u_i \quad (15)$$

We combine these constraints into the Lagrangian \mathcal{L} :

$$\mathcal{L}(u_i) = \lambda(s_i^\top u_i - \delta) + u_i^\top C^{-1} u_i \quad (16)$$

Now, we can derive u_i from this lagrangian, using the constraints $d\mathcal{L}/du_i = d\mathcal{L}/d\lambda = 0$ (derived in 10.2):

$$u_i = \delta \left(\frac{C s_i}{s_i^\top C s_i} \right) \quad (17)$$

How do we choose δ ? For now, we'll keep things simple: if we swap from simple gradient descent to Dan's method, we want ΔJ to be the same in both cases. In other words, we don't want Dan's method to move any more slowly than simple gradient descent.

$$\delta = s_i^\top (-\eta s) \quad (18)$$

4.4.3 Dan's Method Modified

During our work, we realized a potential weakness of this method.

Our original goal was to create an adjustment a_i with covariance C . However, in the above approach, we encourage each adjustment u_i to have covariance C .

At first glance, this shouldn't be too much of a problem: if we're adding many u_i terms with covariance C , then we'll keep the same correlations between different indices.

- However, we know that if we sum independent variables, the covariances should add together.
- In our case, we don't know that each u_i is independent, but the same sort of "gradual accumulation" of covariance could occur, as we add more and more u_i terms together.

The solution is to, instead of encouraging u_i to have covariance C , we encourage the *total adjustment*, $a_i + u_i$, to have covariance C .

We can modify the Mahalanobis distance of Dan's method to accommodate this. Our new lagrangian takes the form

$$\mathcal{L}(u_i) = \lambda(s_i^\top u_i - \delta) + (u_i + a_i)^\top C^{-1}(u_i + a_i) \quad (19)$$

Once again, we derive u_i using the constraints $d\mathcal{L}/du_i = d\mathcal{L}/d\lambda = 0$ (derived in [10.3](#)):

$$u_i = -a_i + (\delta + s_i^\top a_i) \left(\frac{Cs_i}{s_i^\top Cs_i} \right) \quad (20)$$

5 Methods

Now, we get into the details of our implementation and testing of the results.

5.1 Discretizing our state

Our spatial discretization of the ocean breaks it up into a 2D rectangular grid of cells, each having its own temperature. Our code allows for an arbitrary number of rows and columns, but for our experiments, we primarily used a 32 x 32 map.

Our state encodes the "temperature" of the ocean on each of these grid cells. To create our state vector, we stack each column of our 2D map, to create a single column vector x .

Our temporal discretization creates a single state vector state at each evenly-spaced timestep, starting at timestep $t = 0$ and terminating on some arbitrary timestep $t = \tau$.

The state vector at time t is henceforth indicated by $x(t)$.

5.2 Advection-Diffusion-Forcing Model

5.2.1 Continuous Differential Equation

The techniques described above are intended to improve the MITgcm optimization process. However, this model is very time-intensive. Using supercomputers, these models require hours to simulate, and potentially days in order to compute adjoints.

So, instead, to evaluate the effectiveness of these techniques, we chose to use a simplified version of this model. We decided to implement 3 forms of heat transfer: diffusion, advection, and forcing.

$$\frac{\partial x}{\partial t} = \overbrace{K \nabla^2 x}^{\text{Diffusion}} - \overbrace{v \nabla x}^{\text{Advection}} + \overbrace{F(f - x)}^{\text{Forcing}} \quad (21)$$

- F is a scalar coefficient, uniform across our entire map.

5.2.2 Discretization of Differential Equation

This equation was discretized in space and time to first-order, as outlined in [10.4](#).

- ∇x and $\nabla^2 x$ approximately differentiate x "spatially": in our discrete model, this means with respect to the 2D grid cell indices.

Our temporal discretization allows us to approximate $\partial x / \partial t$. With this method, we can "simulate forward" our ocean state, using Euler's method:

$$x(t + 1) \approx x(t) + \left(\frac{\partial x}{\partial t} \right) \Delta t \quad (22)$$

Using a first-order approximation has an additional benefit: it allows us to treat the updated state $x(t + 1)$ as an **affine** function of the current state, x . Thus, we introduce

$$x(t + 1) \approx M \quad (23)$$

5.3 World Parameter Generation

After implementing our advection-diffusion-forcing model, we need a "world" to simulate according to these physics: an initial state, an atmosphere, and additional parameters that impact our physics.

5.3.1 Initial State

Our initial state can be generated randomly, with each grid cell drawn from from a normal distribution.

5.3.2 Atmosphere

A few initial comments:

- In our simplified model, the atmosphere only has a temperature at each grid cell. We don't model wind currents, pushing air (or surface water) between cells.
- Moreover, our atmosphere is *time-invariant*: the atmosphere temperature is the same for all time.
 - Our code can be smoothly adjusted to accommodate for time-variant atmosphere, however.

Now, we can proceed.

We need to generate two version of the atmosphere:

- The **true** atmosphere: this is used to generate the "real" ocean state, which we observe.
- The **first-guess** atmosphere: this is the we want to optimize, to bring it closer to the true atmosphere.

We want a few things from these atmospheres:

- The first-guess atmosphere should be similar to the true atmosphere, but have some inaccurate component.
- Both atmospheres should have similar covariances: we will assume that the source of this first-guess also has access to information about covariance, and will include it in their atmospheric estimate.

Our approach to solve this is to separate these atmospheres into three components:

- f_0 : this is the "known component" of the true atmosphere. This represents the component that our first-guess has accurately predicted: we include it in both atmospheres.
- f_1 : This is the "unknown component" of the true atmosphere. We don't include it in the first-guess.
- f_2 : This is the "error component" of the first-guess atmosphere. It's not part of the true atmosphere: it's the part we get wrong.

To keep the desired covariance structure for the true and first-guess atmospheres, we assume that all three of these terms have covariance C .

- For our simplified case, we compute C using a gaussian dropoff, as described in 10.5.

$$f_0, f_1, f_2 \sim \mathcal{N}(0, C) \quad (24)$$

What percent of the atmosphere does our first-guess atmosphere get right? We represent this with the parameter $\gamma \in (0, 1)$.

With all of these variables, we can finally create our first guess atmosphere f_{guess} , and our true atmosphere f_{true} .

$$f_{\text{true}} = \gamma f_0 + (1 - \gamma) f_1 \quad (25)$$

$$f_{\text{guess}} = \gamma f_0 + (1 - \gamma) f_2 \quad (26)$$

5.3.3 Current Flow

Each grid cell in our ocean should have a current flow v , in order to make use of advection.

- But we need to be careful: we don't want to accidentally end up with a cell where water is flowing in more than it's flowing out.
- Not only is this unrealistic, it also violates the assumptions of our advection-diffusion-forcing differential equation.

Our solution is to create "circulating" current fields. If they flow in loops, whatever water enters a grid cell will come out as it continues along the loop.

We start by creating several "gaussian bumps" on our map: we add together several gaussian functions, randomly placed at different positions on our map.

Then, we take the gradient of this function.

- These gradients will all point towards the top of the nearest hill.

We then rotate these gradients 90 degrees: we replace each vector $[v_x, v_y]$ with $[v_y, -v_x]$.

The result are vectors which flow along the contours of our original plot. These contours typically form "looping" patterns.

5.3.4 Other parameters

In order to make our model more flexible, each of our parameters can be different for each grid cell. For example, two different grid cells might have different x -axis and y -axis side lengths, diffusivities along each edge, areas, etc.

This allows us to model several different types of behaviors: land continents, non-flat planets (since each cell on a globe would need to have a different size and shape), etc.

5.3.5 Cyclic Variables

We also include two additional variables, allowing for our world to be cyclic along the x -axis or the y -axis.

- Allowing east-west cyclic behavior allows us to better model an earth-like planet.

5.4 Observations

Once we have all of our variables, we can simulate the "true" ocean state, and observe it. These observations will be used to optimize f .

At each timestep, we randomly select n grid cells to observe: we choose a different random selection at each timestep.

For every single observation, we add some random gaussian noise to the true state, modelling real-world uncertainty in our observation instruments.

5.4.1 $H_y(t)$ notation

When we were defining our loss function J , why did we represent our observed data as $H_y(t)$?

- We treat $y(t)$ as the full, "true" ocean state, before we observe it.
- H is our **observation operator**: it represents our process of selecting random grid cells, and observing them with noise added.

Thus, $H_y(t)$ is the result of observing our true state.

5.4.2 Unobserved grid cells

How do we deal with unobserved cells? We simply assign them as a NaN value, and ignore them when computing J .

- $H_y(t)$ is the same length as $y(t)$: any cell we don't observe, is left as a NaN value.

Suppose $y_i(t)$ is the i^{th} element of the true state at time t . Meanwhile, $w_i(t)$ is the normally distributed noise (standard deviation σ) we add to it.

$$w_i \sim \mathcal{N}(0, \sigma) \quad (27)$$

$$Hy_i(t) = \begin{cases} y_i(t) + w_i(t) & \text{if observed} \\ \text{NaN} & \text{otherwise} \end{cases} \quad (28)$$

5.5 Computing the gradient

Computing the gradient is, as mentioned previously, handled through the adjoint method (10.7).

5.6 Implementing our gradient descent variants

Each gradient descent variant (code shown in 10.8) uses the same template, only modifying the code that computes u_i between each variation.

5.7 Measuring performance: success metrics

We use three distinct ways of measuring the performance of our gradient descent variants.

5.7.1 Ocean Misfit

This is the loss we're optimizing over, J :

$$\ell_{\text{ocean}} = J(x) = \sum_{t=0}^{\tau} \left(x(t) - Hy(t) \right)^{\top} \left(x(t) - Hy(t) \right) \quad (29)$$

The larger J is, the more inaccurate our ocean simulation is. Our gradient descent procedure is explicitly designed to minimize this.

5.7.2 Atmosphere Misfit

Similar to the "ocean misfit", this is a squared difference:

$$\ell_{\text{atm}} = \left(f_i - f_{\text{true}} \right)^{\top} \left(f_i - f_{\text{true}} \right) \quad (30)$$

This gives us a notion for how different our current estimated atmosphere is from the true atmosphere.

We don't have access to f_{true} during training. Thus, this can be seen as a measure of overfitting: if our ocean misfit is low, but the atmosphere misfit is high, then our procedure isn't actually improving our atmospheric estimate.

- Which is important, since improving our atmospheric estimate is our original goal.

5.7.3 Mahalanobis Distance

Our alternative approaches were all intended to encourage \mathbf{a}_i to have covariance similar to \mathbf{C} .

Given that we measured our covariance similarity using the mahalanobis distance,

$$\ell_C = \mathbf{a}_i^\top \mathbf{C}^{-1} \mathbf{a}_i \quad (31)$$

It is worthwhile for us to investigate how successful each of our four approaches were for this task.

6 Results

6.1 Simple Gradient Descent vs. Modifying J

First, we show off plots of:

- Our simple gradient descent method (3.3)

$$u_i = -\eta s_i \quad (32)$$

- constraining the covariance by modifying J (4.4.1)

$$u_i = -\eta \left(s_i + \alpha \cdot 2C^{-1}a_i \right) \quad (33)$$

Some mild hyperparameter tuning has been applied to α , to get the best results.

6.1.1 Plots




Figure 1: Modifying J seems to slow down our ocean fitting. This make sense, since we're no longer only minimizing J; we have a second term to prioritize, as well.

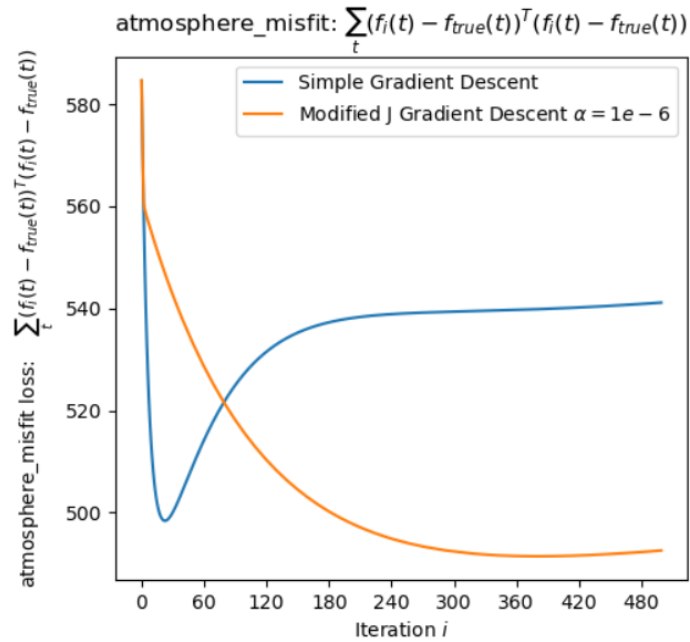


Figure 2: As we might expect, simple gradient descent begins overfitting if we run it for too long.

Modifying J seems to create a better atmospheric estimate!

It also seems to overfit more slowly: this could be because we're using a more realistic adjustment, or because Mahalanobis acts as a regularizer, penalizing larger magnitudes of a_i .

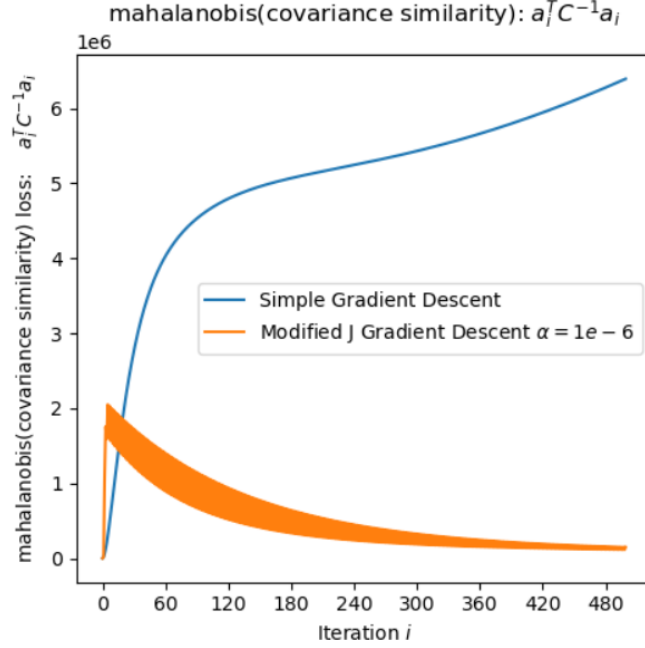


Figure 3: It makes sense that simple gradient descent would have increasing mahalanobis: there's no part of the procedure which penalizes it for creating an adjustment with an unrealistic covariance.

Strangely, the modified gradient descent seems to oscillate very quickly. Still, we can see that it clearly ends up with a much more realistic covariance.

6.1.2 Conclusions

In our case, where the true goal of optimizing the ocean simulation, is to get the best possible atmosphere estimate, modifying J seems to work better than simple gradient descent.

- This makes sense for the same reasons that motivated us to try this approach at all: we're penalizing our atmosphere for having a more unrealistic covariance. So, we might expect that to create a better atmosphere.

That said, it's somewhat unclear why we see such an aggressive oscillation on the Mahalanobis plot.

6.2 Dan Method

Next, we compare our previous methods (3.3 and 4.4.1) to the Dan method (4.4.2).

$$u_i = \delta \left(\frac{C s_i}{s_i^T C s_i} \right) \quad (34)$$

6.2.1 Plots




Figure 4: The Dan method performs almost exactly as well as simple gradient descent! This should make sense, since our optimization includes a constraint that J should decrease by the same amount as it would have normally.

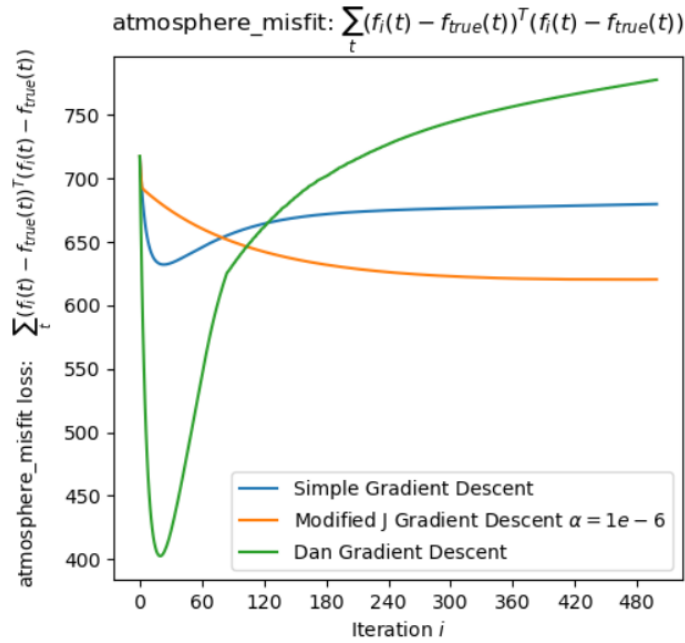


Figure 5: Dan's method creates a much better fit than either other method!

For unknown reasons, however, it suffers more from overfitting. Still, we can remedy this by using an early-stopping condition (stopping once J doesn't meaningfully decrease).

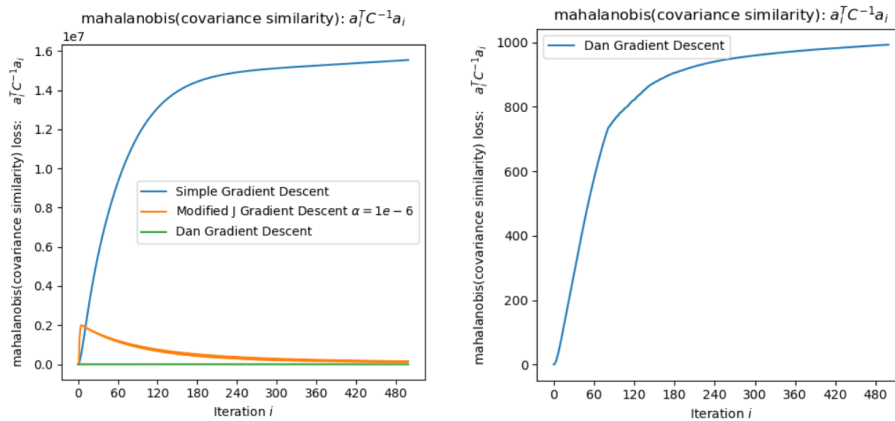


Figure 6: Dan's method keeps the Mahalanobis distance so low, that in our first plot, it's essentially invisible: it presumably does an excellent job at keeping the covariance close to C , suggesting that we have a very realistic adjustment.

The second plot only plots the Dan method, showing that, while Mahalanobis increases much less than the other methods (note that the vertical axis of the left plot is scaled by $1e7$), it does still increase until converging, making a similar curve to simple gradient descent.

We should expect the Mahalanobis to increase from 0 regardless: even if we have the correct covariance structure, we're still increasing the magnitude of our control adjustment from 0.

6.2.2 Conclusions

Dan's method performs best by all of our visible measures, barring some strange overfitting issues.

It outpaces the modified-J approach, so we'll omit that one from further comparisons, for readability.

6.3 Dan Method Modified

Finally, we introduce the Dan Modified Method (4.4.3).

$$u_i = -a_i + (\delta + s_i^\top a_i) \left(\frac{C s_i}{s_i^\top C s_i} \right) \quad (35)$$

6.3.1 Plots

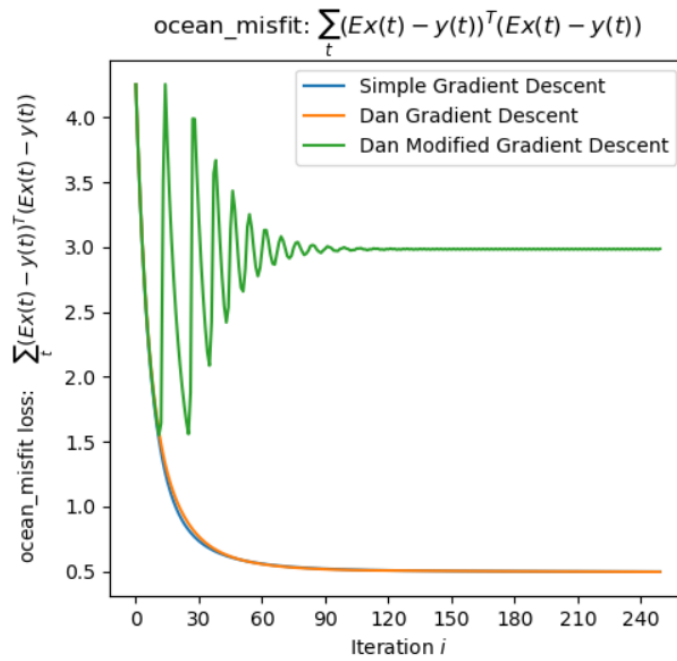


Figure 7: Dan modified starts out performing similarly to the regular Dan method before starting to aggressively oscillate... this is a serious problem.




Figure 8: Same sort of pattern as before: very unstable behavior.

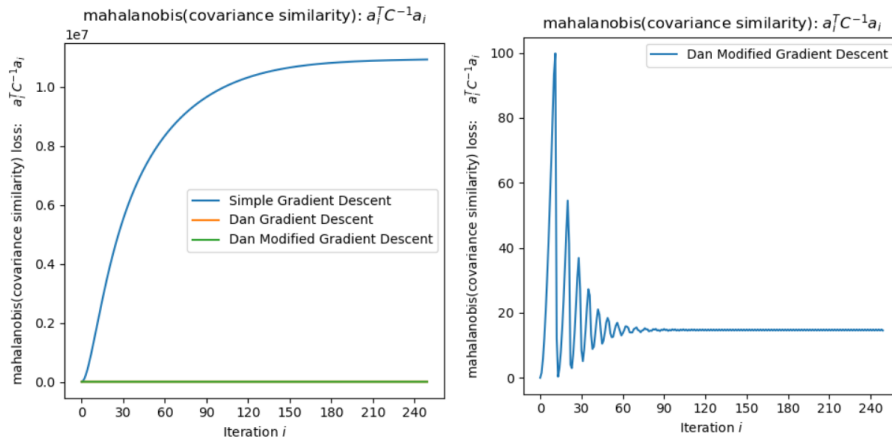


Figure 9: It seems that the mahalanobis distance is getting reset down to nearly 0, over and over again. Just as in the other cases, it seems to eventually converge on some particular mahalanobis value.

6.3.2 Conclusions

The Dan Modified method clearly fails: it starts out performing almost identical to the Dan method, and then abruptly becomes much worse. From there, it starts improving again, only to once again worsen. It repeatedly oscillates, until it converges on a worse result than Dan gradient descent's optimum.

But why does it fail?

6.4 Why does the "Dan Modified" Method Fail?

6.4.1 PCA Analysis

From the above plots, it's not entirely clear *how* the Dan Modified method is failing. Is our control adjustment oscillating over the same values? What happens when it breaks: why is it so abrupt?

The easiest way to test this would be if we could visualize the trajectory across iterations. Initially, this seems infeasible: our atmospheric vector is in a very high-dimensional space (a 32x32 grid becomes a length-1024 vector).

However, it's possible that most of the trajectory behavior is moving in only one or two dimensions: if so, then we can depict our trajectory in those dimensions, without losing information. We don't need to display dimensions that hardly change (have very low "variance"): if they're mostly constant, then we wouldn't learn anything from seeing them.

We can measure this using Principle Component Analysis (PCA): we determine which axes have the greatest variance, and project our data onto those axes.

- Notably, these axes don't have to be aligned with any of the original dimensions in our 1024-dim space: they can instead be a linear combination of our original dimensions.

This is only useful if we're right, and all of the variance can be explained by a few dimensions. As a rule of thumb, we want at least 95% of our variance to be explained by our first two axes, for those to fully describe our trajectory.

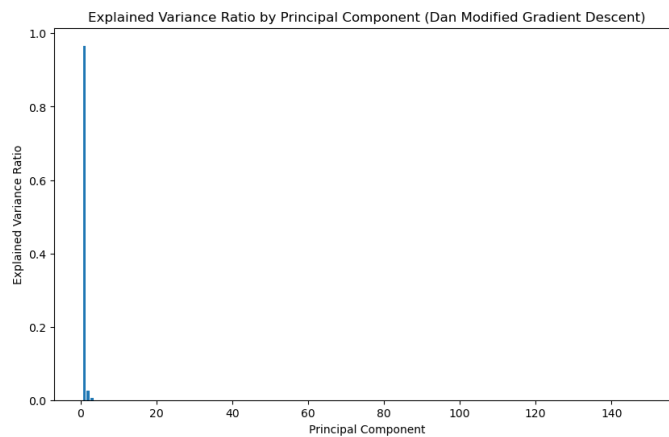


Figure 10: In this plot we see how much variance is explained by each component. Our first component describes 96.63% of the variance, and the second component describes 2.63% of it. So, our first two axes should be sufficient to accurately represent our trajectory.

If that's the case, let's see what result we get:

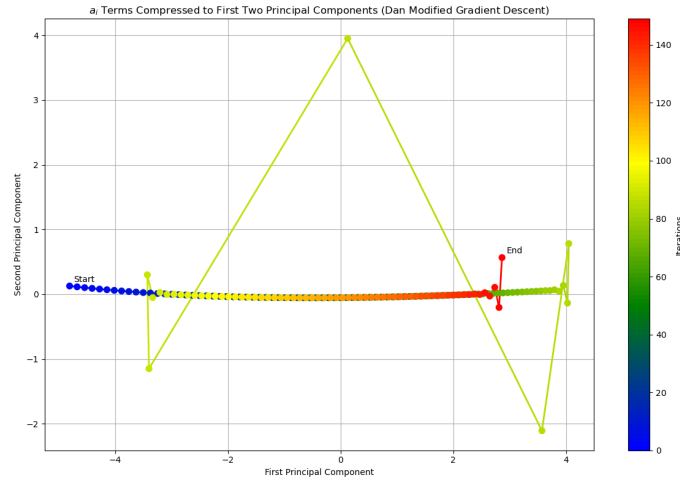


Figure 11: This plot is a bit difficult to read at first. But, if we follow the heatmap showing us the flow in time, it becomes a bit more clear:

Our plot starts out by optimizing, mostly along one axis: moving in, more or less, a straight line (left to right).

Then, it abruptly begins destabilizing, and hopping erratically, mostly perpendicular to the previous path. Once it restabilizes, we end up near where we first started.

Once it restabilizes, it essentially follows the same path that it did before.

We can plot each of our first two principle components to make this behavior clearer:




Figure 12: These plots are a bit easier to read, but they show the same sort of pattern as before: we move in a straight line along PCA 1, until the method destabilizes, hopping along PCA 2, and reversing progress on PCA 1.

Why don't we start at 0, if $a_0 = 0$? Because PCA subtracts the mean of all of our data points: so, our plot has been shifted so the mean of every point is 0.

6.4.2 "Forgetting" theory

In order to see why this method fails, let's return to the solution, from 4.4.2.

$$u_i = -a_i + (\delta + s_i^T a_i) \left(\frac{Cs_i}{s_i^T Cs_i} \right) \quad (36)$$

What happens to our adjustment when we apply this update to it?

$$a_{i+1} = a_i + u_i \quad (37)$$

$$a_{i+1} = a_i - a_i + (\delta + s_i^T a_i) \left(\frac{Cs_i}{s_i^T Cs_i} \right) \quad (38)$$

This method *completely eliminates* our old adjustment, a_i : subtracts it away.

We've essentially "forgotten" our old adjustment: it doesn't contribute to the new position of our adjustment.

$$a_{i+1} = (\delta + s_i^\top a_i) \left(\frac{Cs_i}{s_i^\top Cs_i} \right) \quad (39)$$

Here's how we can interpret this:

- After our previous timestep, we had some adjustment a_i .
- After our new timestep, our adjustment is pointing in a completely new direction: the direction of vector $-Cs_i$ (all other terms in the above equation are scalars, not vectors). Because a_i is completely cancelled out, we move in the $-Cs_i$ direction from the origin.

During our first timesteps, this works fine, because $-Cs_i$ is in the same direction as the previous a_i (as we can see from the previous plots, we're moving in a mostly-consistent direction).

We run into problems if Cs_i starts to change direction after a few timesteps: the further we are from the origin, the more our position changes if we were to rotate slightly. So, we take a very big step from the origin, in some new direction. Thus, the difference between a_i and a_{i+1} is very large, relative to what we expect for a typical gradient descent step.

This is problematic for several reasons:

- The bigger our step is, the less accurate our gradient will be: the gradient *locally* describes the direction of greatest increase.
- Cs_i is computed at the position a_i . But, when we take our new step proportional to $-Cs_i$, we're not starting at position a_i : we're starting from the origin, and then moving in the direction $-Cs_i$.
 - In other words, we're using gradient for position a_i , to determine our direction of motion from the origin.
 - The origin's gradient could very likely be in a different direction: we're using the **wrong gradient** for moving from the origin.

Once we've taken one big, incorrect step, we've moved into a new position, and our gradient will change direction even more. This causes us to take another, bigger step in an even more wrong direction. This repeats, self-amplifying, until the algorithm fails catastrophically.

6.4.3 Supporting evidence: $-Cs_i \cdot a_i$

Of course, to this point, this is only a guess: we need evidence.

This problem hinges on the idea that when the our algorithm breaks, it occurs because \mathbf{a}_i and $-\mathbf{Cs}_i$ are in different directions. This, of course, is only possible if \mathbf{a}_i and $-\mathbf{Cs}_i$ are in different directions at the point where our function breaks, and the same direction otherwise.

We can measure this by taking the normalized dot product between them:

$$\frac{-\mathbf{Cs}_i \cdot \mathbf{a}_i}{|\mathbf{Cs}_i||\mathbf{a}_i|} \quad (40)$$

The dot product is 1 if two vectors are in the same direction, 0 if they're perpendicular, and -1 if they're in directly opposing directions.

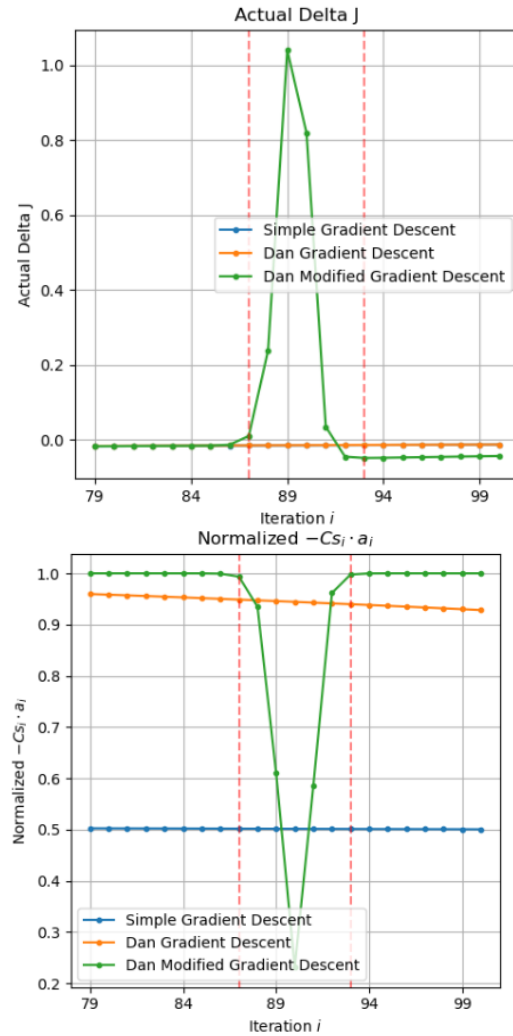


Figure 13: We can use the true ΔJ to see when our algorithm breaks: we're no longer successfully improving our ocean misfit; we're making it worse.

Notably, over the same region, we see exactly what we expected: whereas $-Cs_i$ and a_i were previously in the same direction, we start to run into problems once this is no longer the case.

We use red vertical lines to indicate this "failure region" of the Dan Modified method.

6.4.4 PCA for comparing methods

An alternative explanation might be that Dan Modified happens to create adjustments that, eventually, enter some unusually unstable region of the adjustment space.

The first way we could check this is to compare the adjustments of the Dan Modified Method, to those of the other methods. How similar/different are they?

We can visualize this, once again, by using PCA. We combine all of the adjustments from all three of our methods (simple gradient, Dan, Dan Modified) into a single dataset, and

take PCA over it.

When we compute the variance explained by each dimension, we once again find that our first two components describe $> 95\%$.

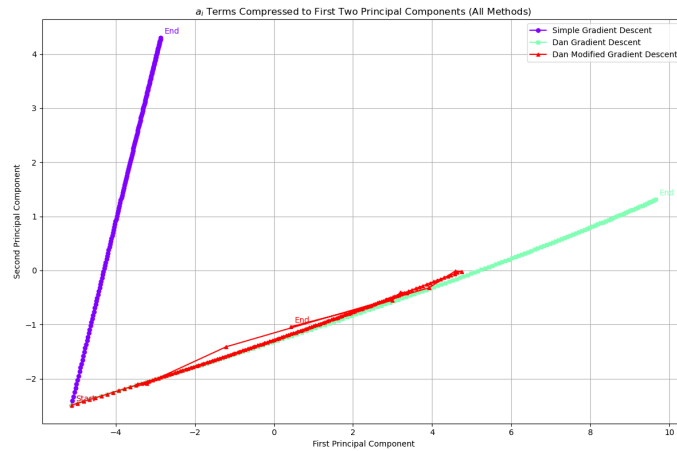


Figure 14: Our first two PCA components explain 83.5% and 15.2% of the variance, respectively.

When we compare these plots, we see that, at least visual inspection, the Dan method and Dan Modified method seem to take very similar trajectories.

Notably, simple gradient descent takes a much more different trajectory.

We're not done yet, though.

6.4.5 Switching Methods

There's a more sure way that we can test this theory. If Dan Modified is breaking because of the region of the adjustment space, then we should be able to place the Dan method and Dan Modified method in the exact same spot, and either they both break, or neither one breaks.

First, we'll test this by running the Dan modified method, and noting where it breaks. If we start from the adjustment α_i where the Dan Modified method breaks, and switch to the Dan method, then we can see whether it still breaks. If it doesn't, then the method, not the adjustment α_i , is the problem: they'll be using the exact same prior adjustment.

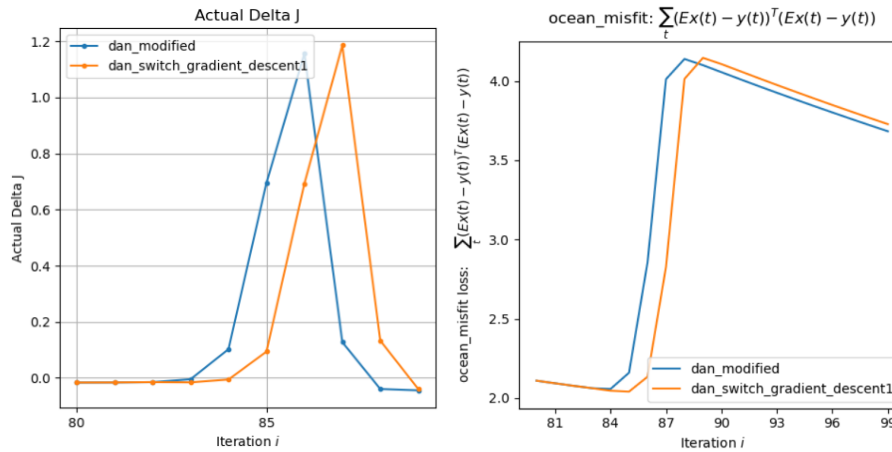


Figure 15: In the "danswitch" method, we used the normal Dan method for one timestep. During this timestep, our optimization proceeded as normal: it didn't break until we switched back to the Modified Dan method on the next timestep.

What if we use the Dan method for longer? Maybe it breaks, but it takes longer?

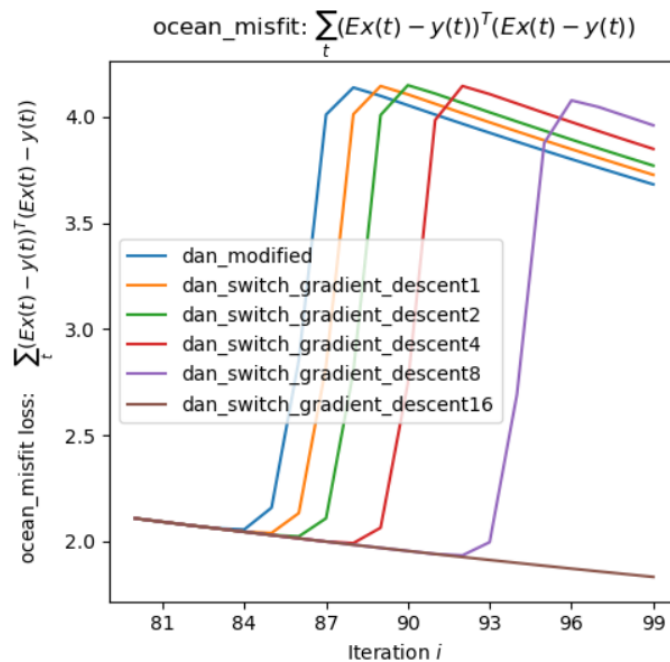


Figure 16: Each method is labelled with the number of iterations for which we switch to the Dan method. It seems that the Dan method doesn't break at all, but if we switch back to the Dan modified method, it'll break immediately.

This seems to support our theory. Another test would be to go the other way: start with the Dan Method, and switch to the Dan Modified method.




Figure 17: The results here are interesting: if we switch from Dan to Dan Modified, we can actually break our algorithm *faster* than if we just used the Modified Dan method. Why might that be?

Well, we're assuming that Dan Modified breaks because of an issue with $-Cs_i$ and a_i pointing in different directions.

So, let's check the dot product between the two:

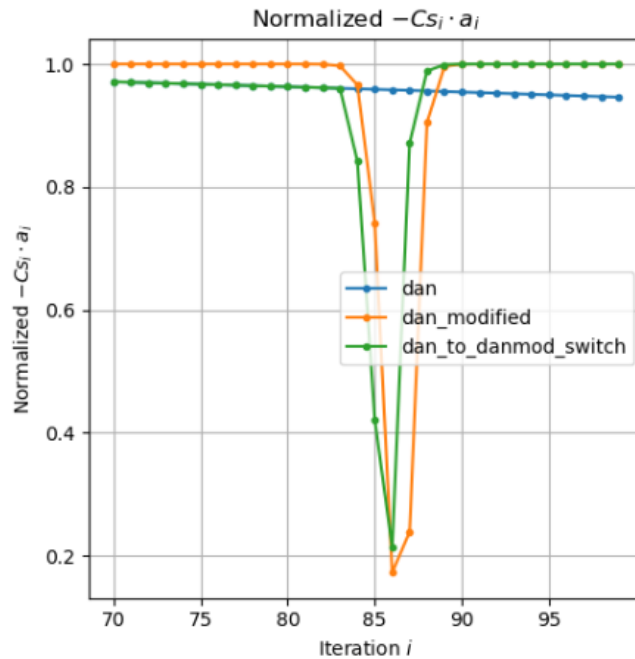


Figure 18: It seems that the simple Dan method doesn't generally have $-Cs_i$ and a_i pointing in the same direction: the normalized dot product is a bit less than 1.

This makes sense: the reason the Modified Dan Method has $-Cs_i$ and a_i pointing in the same direction, is that after each iteration, we completely replace a_i with a vector pointing towards $-Cs_i$. The same isn't true for the simple Dan method.

If we trust our theory for why Dan Modified breaks, this also explains the above behavior. If the Dan Modified method breaks because a_i and $-Cs$ are pointed in different directions, and the simple Dan method develops this property more quickly, then it would make sense for it to break when we switch to Modified Dan.

6.5 Is the Dan method really flawed at all?

It seems that the Modified Dan method is a dead end; even if we're incorrect about the reason why it breaks, we still know that it consistently does.

So, why did we decided to try the Modified Dan method? Wasn't there a problem with the simple Dan method? But when we run the simple Dan method, it seems to run well. In fact, the mahalanobis distance doesn't continuously increase like we thought it might: it actually begins to plateau.

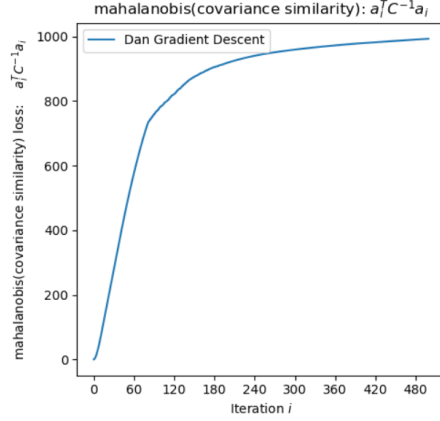


Figure 19: A reminder of how our Dan Method's mahalanobis distance increases.

6.5.1 Mahalanobis pressures Correlation, not Covariance

Eventually, we came to realize the mistake in our thinking: let's review the lagrangian for the Dan method (4.4.2).

$$\mathcal{L}(u_i) = \lambda(s_i^\top u_i - \delta) + u_i^\top C^{-1} u_i \quad (41)$$

Here, we are trying to minimize the mahalanobis distance. However, we misunderstood the purpose of the mahalanobis distance:

- At first, we thought of it as constraining u_i to have covariance C .
- *But*, Mahalanobis has a regularization effect as well: if we decrease the magnitude of u_i , that will also decrease the mahalanobis.

Suppose we have a vector z with covariance C . The vector $z/2$ would have a lower mahalanobis distance, even though the covariance is now $C/2$.

- In other words, the Mahalanobis distance pressures our covariance to be *proportional* to C , but it wants the vector magnitude to be as small as possible.

So, it encourages the "structure" that covariance C has, regardless of magnitude. We're not enforcing covariance: we're enforcing the **correlation** of our vector.

- Vector z_1 and z_2 having the same correlation is equivalent to having a *proportional* covariance: $V[z_1] = rV[z_2]$.

Why is this significant? Let's think of our previous concern:

- We were worried that each update u_i would have covariance C . As we add many updates, a_i will end up with a covariance much larger than C : it would have the right *correlation*, but it would be too large.

But this isn't how the mahalanobis constraint works: it pressures u_i to have a same covariance as implied by C . The covariance of u_i can be much smaller than C .

In fact, the above plot suggests that, as we run our gradient descent, u_i makes smaller and smaller adjustments to the covariance with later iterations. It seems to converge on some constant mahalanobis.

So, it seems that, while the Dan Modified method may not be functional, the Dan Method doesn't have the problem we thought it did.

Since each u_i should have roughly the right correlation, and thus a covariance *proportional* to C , then we should be adding to get a total adjustment a_i with covariance proportional to C , as well.

But does a_i converge on covariance C ? Or some multiple of it? Let's investigate.

6.5.2 The Mahalanobis Distance we expect

We've been measuring our covariance similarity to C with the Mahalanobis distance. But as we just established, the Mahalanobis distance also penalizes magnitude:

- Having a mahalanobis distance of 0 doesn't mean you have the desired covariance: it means that $a_i = 0$.

So then, how do we know that our covariance is correct? Well, we need to find out what the mahalanobis would be, for a vector with the correct covariance.

Testing this is simple: we generate some random normal vector z with covariance C , and measure their mahalanobis.

- For our length-1024 vector, we find that the mahalanobis is roughly 1000.
- Is it the case that for a vector of length n , the average mahalanobis is n ?
- Upon testing other examples, the answer is yes!

The reason why this occurs is discussed in [10.6](#).

More importantly, we can compare this "typical" mahalanobis value for a length-1024 vector, to the mahalanobis that the Dan Method converges to.

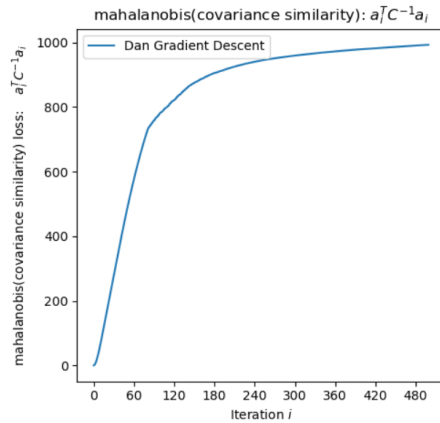


Figure 20: The Dan method converges to ≈ 1000 , too!

This suggests that the Dan method likely converges on the correct covariance!

However, we can't be sure: Mahalanobis doesn't allow us to clearly separate "covariance similarity" from "magnitude".

- If we compare their magnitudes, the "typical" random length-1024 vector with covariance C (covariance matrix defined by 10.5) has magnitude ≈ 17 , while our final a_i has magnitude ≈ 14 .
- So, they have similar magnitudes: we could investigate further, but for now, we'll take this as sufficient evidence that the covariance of a_i (at least approximately) converges to the desired covariance C .

6.5.3 Conclusion

Our takeaways:

- The modified-J method appears to work better than simple gradient descent, but worse than Dan's Method.
- Dan's Method works best, and appears to cause our adjustment to approach the desired covariance.
- Dan's Modified Method does not work as intended, and is unlikely unnecessary.

7 Future Work

7.1 Extensions of our Toy Model

- Our atmosphere is current time-invariant. We could model a time-varying atmosphere.
- Our ocean currents are time-invariant as well, presently.
- We could test more loss functions, or try to figure out why J-modify doesn't perform as well as Dan's Method.
- Introduce more parameters: wind currents, for example.
- Our model was entirely ocean, on a flat grid, with uniform grid cell dimensions. We could introduce land masses, or reshape our model into a more realistic topology (like a low-resolution globe).
- We could advance our model beyond a first-order estimate.
- In the same way that we used covariance constraints on the atmosphere, we could introduce information about the covariance of the ocean state. This would allow us to infer about ocean states we do not observe.
- Our code is capable of using a separate covariance for the control, versus the control adjustment. We could see how this affects our performance.

7.2 Extensions Beyond our Toy Model

- The original goal of testing Dan's Method was to apply it to MITgcm. Dan's method could be implemented in that environment.
- Our model uses gradient descent, but we could use other optimization approaches more similar to MITgcm, like the BFGS algorithm.
- Our covariance matrix is manufactured: an extension should look into finding a more accurate covariance matrix, based on real data.
 - However, for a huge model like MITgcm, it would be very difficult to store this enormous matrix. So, we would need to find more space-efficient ways to store/compress this matrix.
 - One option is SVD, where we only retain the first n singular values.

7.3 Further investigation on our work

- When Dan's modified method fails...
 - How does this behavior relate to the eigenvectors of C ? Can we characterize the failure based on projecting our adjustment onto the eigenvectors of C ?

- Why does it consistently loop over the same region of space as it repeatedly fails? And why does it converge to a point in the middle of that region?
- Are there other ways to test our theory for why it fails?

8 References

- [1] Dan Amrhein. Ocean state estimation with atmospheric adjustments constrained by prior covariance and observations. Unpublished manuscript, April 2024.
- [2] Andrew M. Bradley. Pde-constrained optimization and the adjoint method. https://cs.stanford.edu/~ambrad/adjoint_tutorial.pdf, Oct 2019.
- [3] Shumon Koga. Report: Project in jpl. Technical report, Jet Propulsion Laboratory (JPL), November 2017. Unpublished report.

9 Acknowledgments

I thank my mentor Ian Fenty, co-mentor Ou Wang, and external mentor Dan Amrhein.

I would also like to thank Caltech for providing resources through SFP, and NASA for providing a workspace, work materials, site access, etc.

I studied the adjoint with the help of several valuable resources, including [\[2\]](#), [\[3\]](#).

10 Appendices

10.1 Matrix Conventions

In this paper, we assume that vectors default to column-vector form.

$$z = \begin{bmatrix} z_1 \\ z_2 \\ \vdots \\ z_n \end{bmatrix} \quad (42)$$

Moreover, we use **denominator layout** for matrix derivatives take x and y to be scalars, and \vec{x} and \vec{y} to be vectors.

$$\frac{\partial y}{\partial \vec{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix} \quad (43)$$

$$\frac{\partial \vec{y}}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x} & \frac{\partial y_2}{\partial x} & \dots & \frac{\partial y_m}{\partial x} \end{bmatrix} \quad (44)$$

$$\frac{\partial \vec{y}}{\partial \vec{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_2}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_1} \\ \frac{\partial y_1}{\partial x_2} & \frac{\partial y_2}{\partial x_2} & \dots & \frac{\partial y_m}{\partial x_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_n} & \frac{\partial y_2}{\partial x_n} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix} \quad (45)$$

10.2 Dan's Method Derivation

10.2.1 Constraints of Dan's Method

Dan's method can be interpreted as a modification to gradient descent.

- In its original writing, it appears to only reference a **single iteration**.

More accurately, it only discusses the total control adjustment u .

In particular, we modify our **gradient update**, u_i , with **two constraints**.

However, it only uses a single sensitivity s : since s is re-computed each GD iteration, it's most natural to treat this as a single iteration of GD.

First, we want to ensure that u_i productively improves our loss function.

Constraint 1

The **total change** in J over this iteration is equal to some constant δ .

- We'll use the approximation $\Delta J \approx u_i^\top s$.

$$u_i^\top s = \delta \quad (46)$$

But our real goal is to encourage u_i to have an expected covariance C :

The following technique works under the assumption that u_i has mean 0. We'll make this assumption.

Constraint 2

We want to make the **covariance** of u_i as close to C as possible.

- We measure this with the **mahalanobis distance**

$$u_i^\top C^{-1} u_i \quad (47)$$

- The larger it is, the further we are from having covariance C .

So, we want to **minimize** this distance.

We can combine these into a single loss function:

Key Equation 3

Our **loss function** \mathcal{L} for finding our update u_i is given by:

$$\mathcal{L}(u_i) = 2\lambda(u_i^\top s - \delta) + u_i^\top C^{-1} u_i \quad (48)$$

This encourages u_i to have covariance C .

10.2.2 Optimizing over our constraints

How do we use this loss function from 48? We set two constraints: $d\mathcal{L}/du_i = 0$, and $d\mathcal{L}/d\lambda = 0$.

Because our derivatives equal 0, we divided by 2 before getting the below equations.

$$0 = \frac{d\mathcal{L}}{du_i} = \lambda s + C^{-1}u_i \quad (49)$$

$$0 = \frac{d\mathcal{L}}{d\lambda} = u_i^\top s - \delta \quad (50)$$

For our derivation below, we'll find the rearranged version more useful:

$$-C^{-1}u_i = \lambda s \quad (51)$$

$$\delta = u_i^\top s \quad (52)$$

10.2.3 Derivation

From here, we do some algebra.

From 51, we find

$$-C^{-1}u_i = \lambda s \quad \Rightarrow \quad u_i = -C\lambda s$$

$$u_i = -C\lambda s \quad (53)$$

Plugging 53 into 52, we find:

$$\delta = u_i^\top s \quad \Rightarrow \quad \delta = (-C\lambda s)^\top s = -\lambda s^\top C^\top s = -\lambda s^\top C s \quad \Rightarrow \quad \delta = -\lambda s^\top C s \quad \Rightarrow \quad \lambda = \frac{-\delta}{s^\top C s}$$

$$\lambda = \frac{-\delta}{s^\top C s} \quad (54)$$

Plugging 54 into 53, we get:

$$u_i = -C\lambda s \quad \Rightarrow \quad u_i = -C\left(\frac{-\delta}{s^\top C s}\right)s \quad \Rightarrow \quad u_i = \delta\left(\frac{Cs}{s^\top C s}\right)$$

This gives us our gradient update u_i :

Key Equation 4

According to **Dan's method**, we have computed our gradient update:

$$\mathbf{u}_i = \delta \left(\frac{\mathbf{C}\mathbf{s}}{\mathbf{s}^\top \mathbf{C}\mathbf{s}} \right) \quad (55)$$

10.3 Dan's Modified Method Derivation

Based on constraining α_{i+1} 's covariance instead of u_i , let's adjust our previous equations:

- We still want our **newest** update, u_i , to create a fixed change in J , δ .

$$u_i^\top s = \delta \quad (56)$$

- However, we want to encourage our **total** control adjustment, $\alpha_i + u_i$, to have covariance C .

$$(\alpha_i + u_i)^\top C^{-1}(\alpha_i + u_i) \quad (57)$$

Constraint 5

We want to encourage the covariance of **control adjustment** $\alpha_i + u_i$ to be close to C .

- We'll do this by minimizing the **mahalanobis distance**

$$(\alpha_i + u_i)^\top C^{-1}(\alpha_i + u_i) \quad (58)$$

We'll make a new loss function:

Key Equation 6

Our **lagrangian** \mathcal{L} for finding our update u is given by:

$$\mathcal{L}(u_i) = 2\lambda(u_i^\top s - \delta) + (\alpha_i + u_i)^\top C^{-1}(\alpha_i + u_i) \quad (59)$$

This encourages $\alpha_i + u_i$ to have covariance C .

10.3.1 Optimizing over constraints

We differentiate 59 just as we did before:

$$0 = \frac{d\mathcal{L}}{du_i} = \lambda s + C^{-1}(\alpha_i + u_i) \quad (60)$$

$$0 = \frac{d\mathcal{L}}{d\lambda} = u_i^\top s - \delta \quad (61)$$

If we rearrange for utility, we get:

$$-C^{-1}(a_i + u_i) = \lambda s \quad (62)$$

$$\delta = u_i^\top s \quad (63)$$

10.3.2 Derivation

Once again, we manipulate some algebra.

From 62, we find:

$$-C^{-1}(a_i + u_i) = \lambda s \quad \Rightarrow \quad a_i + u_i = -C\lambda s \quad \Rightarrow \quad u_i = -a_i - \lambda Cs$$

$$u_i = -a_i - \lambda Cs \quad (64)$$

Plugging 64 into 63, we find

$$\delta = u_i^\top s \quad \Rightarrow \quad \delta = (-a_i - \lambda Cs)^\top s = -a_i^\top s - \lambda s^\top C^\top s = -(a_i^\top + \lambda s^\top C) s = -(a_i^\top s + \lambda s^\top Cs)$$

We can rearrange for λ :

$$-\delta = a_i^\top s + \lambda s^\top Cs \quad \Rightarrow \quad -(\delta + a_i^\top s) = \lambda s^\top Cs \quad \Rightarrow \quad \lambda = -\left(\frac{\delta + a_i^\top s}{s^\top Cs}\right)$$

$$\lambda = -\left(\frac{\delta + a_i^\top s}{s^\top Cs}\right) \quad (65)$$

Plugging 65 into 64, we get:

$$u_i = -a_i - \lambda Cs \quad \Rightarrow \quad u_i = -a_i - C\left(\frac{-\delta - a_i^\top s}{s^\top Cs}\right)s \quad \Rightarrow \quad u_i = -a_i + \left(\delta + a_i^\top s\right)\left(\frac{Cs}{s^\top Cs}\right)$$

We have our gradient update u_i :

Key Equation 7

According to **Dan's modified method**, we have computed our gradient update:

$$u_i = -a_i + \left(\delta + a_i^\top s\right)\left(\frac{Cs}{s^\top Cs}\right) \quad (66)$$

10.4 The Advection-Diffusion-Forcing Model, Discretized

Here, we briefly outline the discretization of our differential equation.

We simplify this by representing it in 1D. The full, 2D representation can be seen in the code.

Let c_i be the temperature of grid cell i .

10.4.1 Diffusion

We discretize the differential equation

$$\frac{\partial c}{\partial t} = K \nabla^2 c \quad (67)$$

Giving us

$$\frac{\Delta c_i}{\Delta t} \approx K \left(\frac{c_{i+1} + c_{i-1} - 2c_i}{(\Delta x)^2} \right)$$

Note that Δx is the horizontal distance between two cells, or equivalently, the length of one cell.

10.4.2 Forcing

The forcing term

$$\frac{\partial c}{\partial t} = F(f - c) \quad (68)$$

Doesn't require any discretization, beyond discretizing f_i and c_i in space.

$$\frac{\Delta c_i}{\Delta t} = F(f_i - c_i)$$

10.4.3 Advection

The advection expression

$$\frac{\partial c}{\partial t} = -v \nabla c \quad (69)$$

Has to be discretized a bit unusually, in order to avoid some numerical effects.

- If the current pushes right (positive), then heat only comes from the left; we don't care about temperature on the right.
- Vice versa if the current pushes left.

$$\frac{\Delta c_i}{\Delta t} \approx \left(\max(v_i, 0) c_{i-1} + \min(v_i, 0) c_i - \max(v_{i+1}, 0) c_i - \min(v_{i+1}, 0) c_{i+1} \right) / \Delta x$$

This expression may seem nonlinear, but because we assume our current to be time-independent, we can pre-compute each min and max, and turn it into a simple constant.

10.4.4 Affine Model

When we add all three of these expressions together, we can rearrange them to give us an affine model. Using matrix M and constant F , we get the function:

$$c(t + 1) \approx Mc(t) + Ff \quad (70)$$

Which we use for both simulation and adjoint computation.

10.5 Our "gaussian decay" covariance assumption

We need to choose a particular covariance matrix C to work with. We chose a simple approximation: nearby cells should have a greater covariance.

We model this with a gaussian function.

$$C_{a,b} = \sigma_{a,b}^2 = e^{-||a-b||^2/2s} \quad (71)$$

We combine all of these into our covariance matrix:

$$C = \begin{bmatrix} \sigma_{1,1}^2 & \cdots & \sigma_{1,n}^2 \\ \vdots & \ddots & \vdots \\ \sigma_{n,1}^2 & \cdots & \sigma_{n,n}^2 \end{bmatrix} \quad (72)$$

10.6 The χ -squared distribution

Earlier, in 6.5.2, we noticed that, if a length- n vector has covariance C , its expected mahalanobis distance seems to be n . Why is that?

10.6.1 Informal Justification

First, we'll informally justify it. Let's start by consider what the mahalanobis distance *is*. For review:

$$z^T C^{-1} z \quad (73)$$

This equation is a generalized version the 1d equation, where we compute z^2 , and divide by the variance.

$$\left(\frac{z}{\sigma}\right)^2 \quad (74)$$

Since z is assumed to have mean 0, z/σ is a standard normal variable, $z/\sigma \sim \mathcal{N}(0, 1)$.

We're *squaring a standard normal variable*: this is closely related to the definition of the **chi-squared distribution**.

In particular, χ_k^2 is the sum of k squared standard normal variables.

Now, we remind ourselves that z is not actually a single real number: it's a vector of length n .

To make this informal argument easier, let's trust that C^{-1} "cancels out" the covariance C of the random normal vector z . So, we end up with w , a standard normal vector.

$$z^T C^{-1} z = w^T w = \sum_i w_i^2 \quad (75)$$

Each w_i is now a standard normal vector. So, we're summing n squared standard normal vectors: we have a chi-squared distribution, χ_n^2 .

The expected value of this distribution is n : this is exactly what we observe.

10.6.2 Formally demonstrating this result

Our goal is to formally justify the step,

$$z^T C^{-1} z = w^T w \quad (76)$$

Above, we used the "square root" of variance, σ , to cancel out the variance of z in the 1D case. Here, we'll do something similar, and define a sort of "square root" for our matrix C .

Thus we introduce the Cholesky Decomposition:

- If a matrix is positive-definite symmetric (like any covariance matrix C), then it has a "cholesky decomposition into a matrix L , where

$$LL^T = C \quad (77)$$

For this proof, we do not care about the actual contents of L : we just need to know that it exists.

An important property of the cholesky decomposition (or any other "square root" of C):

- **Lemma:** if $w \sim \mathcal{N}(0, I)$, then $Lw \sim \mathcal{N}(0, C)$.

Since we know that Lw is mean zero, we can get the covariance as:

$$\text{Cov}(Lw) = E[(Lw)(Lw)^T] \quad (78)$$

If we rearrange:

$$E[(Lw)(Lw)^T] = E[Lww^T L^T] = LE[ww^T]L^T \quad (79)$$

We already know that w has covariance I . Since it has mean 0, $E[ww^T] = \text{Cov}(w) = I$.

$$\text{Cov}(Lw) = LE[ww^T]L^T = LIL^T = LL^T = C \quad (80)$$

Thus, $Lw \sim \mathcal{N}(0, C)$.

Now, we can use a trick: we know that $z \sim (0, C)$ has the same distribution as Lw . So, we'll substitute one for the other: $z^T C^{-1} z$ is equivalent in distribution to $(Lw)^T C^{-1} (Lw)$.

$$(Lw)^T C^{-1} (Lw) = w^T L^T C^{-1} Lw \quad (81)$$

We'll use the fact that $C^{-1} = (LL^T)^{-1} = (L^T)^{-1} L^{-1}$.

$$w^T L^T C^{-1} Lw = w^T L^T (L^T)^{-1} L^{-1} Lw = w^T w \quad (82)$$

We've shown that our mahalanobis distance is equivalent in distribution to $w^T w$: in other words, equivalent in distribution to χ_n^2 . Proof complete.

10.7 The Adjoint Method

The adjoint method allows us to compute dJ/df more efficiently. Here, we derive/justify the approach.

10.7.1 Setting Up/Motivation

For our derivation, it's easiest to focus on the atmosphere at a single timestep, even though our atmosphere is uniform. We choose an arbitrary timestep q .

Our goal is to modify our atmospheric forcing $f(q)$ to improve our simulation (in other words, reducing J). This can be best represented by asking, "how does modifying $f(q)$ affect J ?" This question is answered by the derivative,

$$\frac{dJ}{df(q)}$$

We can use this to directly compute an adjustment to $f(q)$, to improve our estimate. So, this derivative is our goal.

10.7.2 Our Model

How does $f(q)$ affect J ? It doesn't directly show up in the equation for J .

- Rather, it *indirectly* affects J , by modifying the (simulated) ocean state, $x(t)$.

This effect is represented by our equation for simulating forward in time:

$$x(t+1) = Mx(t) + Ff(t)$$

$f(q)$ influences the next state $x(q+1)$, which contributes to J . But, we're forgetting a second way that $f(q)$ can affect J : by affecting *future states*.

- While $f(q)$ only directly affects $x(q+1)$, we use $x(q+1)$ to compute $x(q+2)$. We can then use $x(q+2)$ to compute $x(q+3)$, and so on.
- So, $f(q)$ affects all of our future states!
- By affecting each of these states, $f(q)$ can affect J at $\tau - q$ different states.

We can account for all of these terms using the multivariable chain rule:

$$\frac{dJ}{df(q)} = \sum_{t=q+1}^{\tau} \frac{dx(q+1)}{df(q)} \cdot \frac{dx(t)}{dx(q+1)} \cdot \frac{\partial J}{\partial x(t)}$$

We know how to compute each of these terms: the first and third terms are known matrix derivatives, so we'll put them off until later.

It's useful to think of this in a second way: above, we've listed every way that $x(q+1)$ can affect J . We have a *total derivative* of J with respect to $x(q+1)$.

$$\frac{dJ}{df(q)} = \frac{dx(q+1)}{df(q)} \left(\sum_{t=q+1}^{\tau} \frac{dx(t)}{dx(q+1)} \cdot \frac{\partial J}{\partial x(t)} \right) = \frac{dx(q+1)}{df(q)} \left(\frac{dJ}{dx(q+1)} \right)$$

10.7.3 Redundant Calculations

This technique gets the job done, but it can be inefficient to use for multiple timesteps: we have a lot of duplicate calculations. Consider an example:

- $f(1)$ and $f(2)$ both affect $x(3)$, which in turn affects J . Thus, both equations require $\frac{dJ}{dx(3)}$.

$$\begin{aligned} \frac{dJ}{df(1)} &= \frac{dx(2)}{df(1)} \left(\overbrace{\frac{dJ}{dx(2)}}^{\text{Total effect of } x(2)} \right) = \frac{dx(2)}{df(1)} \left(\overbrace{\frac{\partial J}{\partial x(2)}}^{x(2) \text{ effect by itself}} + \overbrace{\frac{dx(3)}{dx(2)} \frac{dJ}{dx(3)}}^{x(2) \text{ effect via future timesteps}} \right) \\ \frac{dJ}{df(2)} &= \frac{dx(3)}{df(1)} \left(\frac{dJ}{dx(3)} \right) \end{aligned}$$

10.7.4 The Adjoint Method: Base Case

It seems that, in the above case, it would make sense to compute $dJ/dx(3)$ first, so we can re-use it for computing $dJ/dx(2)$.

- But if we just showed that $dJ/dx(3)$ is used for twice, doesn't it make sense that the same is true for $dJ/dx(4)$?
 - If we use an identical argument to before, we could show that computing $dJ/dx(3)$ involves computing $dJ/dx(4)$.
 - So, we should handle $dJ/dx(4)$ first.

We can use the same logic over and over, going further forward in time: it seems we're reusing a lot of calculations!

The natural conclusion is for us to start with the very last timestep, $dJ/dx(\tau)$.

- Because there are no future timesteps, $x(\tau)$ can only affect J directly:

$$\frac{dJ}{dx(\tau)} = \frac{\partial J}{\partial x(\tau)}$$

10.7.5 The Adjoint Method: Recursion

Now, we can move one step **backwards** in time, using the equation we wrote above:

$$\frac{dJ}{dx(\tau-1)} = \overbrace{\frac{\partial J}{\partial x(\tau-1)}}^{\text{x}(\tau-1) \text{ effect by itself}} + \overbrace{\frac{dx(\tau)}{dx(\tau-1)} \frac{dJ}{dx(\tau)}}^{\text{x}(\tau-1) \text{ effect via } x(\tau)}$$

To make things clearer, we'll rename the variable we're recursively building up:

$$\lambda_t = \frac{dJ}{dx(t)}$$

Rewriting our equation:

$$\lambda_{\tau-1} = \frac{\partial J}{\partial x(\tau-1)} + \frac{dx(\tau)}{dx(\tau-1)} \lambda_{\tau}$$

We get something that looks like a **recursive** relation: $\lambda_{\tau-1}$ references the next element in the sequence, λ_{τ} . As we move further back in time, we find the exact same equation, confirming our suspicions. If we write it in general, we get:

$$\lambda_t = \begin{cases} \frac{\partial J}{\partial x(t)} + \frac{dx(t+1)}{dx(t)} \lambda_{t+1} & \text{if } t < \tau \\ \frac{\partial J}{\partial x(t)} & t = \tau \end{cases}$$

These are our **adjoint variables**.

10.7.6 Using the adjoint

We can find our adjoint variables by moving backwards in time: we start by computing λ_{τ} , and begin decrementing through $t = \tau - 1, \tau - 2, \dots, 2, 1$.

Once we've finished, it's easy to compute our final derivatives:

$$\frac{dJ}{df(q)} = \frac{dx(q+1)}{df(q)} \lambda_{q+1}$$

If we apply this to our model from 10.4 ($x(t+1) = Mx(t) + Ff(t)$), we find that $\frac{dx(t+1)}{dx(t)} = M^T$

$$\lambda_t = \begin{cases} \frac{\partial J}{\partial x(t)} + M^T \lambda_{t+1} & \text{if } t < \tau \\ \frac{\partial J}{\partial x(t)} & t = \tau \end{cases}$$

If we work through the induction, we can simplify this to:

$$\lambda_k = \sum_{i=k}^{\tau} \left((M^{i-k})^T \frac{\partial J}{\partial x(i)} \right)$$

And finally:

$$\frac{dJ}{df(q)} = F\lambda_{q+1} = F \sum_{i=q+1}^{\tau} \left((M^{i-q-1})^T \frac{\partial J}{\partial x(i)} \right)$$

Notice that the last forcing, $f(\tau)$, actually has no effect on our loss: it would be applied to a future state $x(\tau + 1)$, that doesn't exist.

- In the above equation, this would refer to some non-existent $\lambda_{\tau+1}$.

10.7.7 Why is the adjoint useful?

Something worth addressing:

Q: *Couldn't we have computed the answer in our original form, without invoking the adjoint? We could've just plugged values into the chain rule we started with.*

In this particular case, this is true. However, this is only simple, because our model takes on such a simple form, where we can multiply by A^T repeatedly to get our answer.

In many situations, our model can be too complex to get an analytical derivative. So, instead, we might use a more demanding approach, like **finite difference approximation**:

- Modify one variable of $f(q)$ and simulate the whole model, seeing how the loss changes.
- We repeat this process for each variable in $f(q)$, to get the overall derivative.
- Then, we have to repeat *all* of that, for every timestep q .

Using the adjoint method, we can significantly cut down on the work we have to do:

- First, we compute the adjoint variables λ_t : this requires computing our derivatives $\partial J / \partial x(t)$ and $\partial x(t+1) / \partial x(t)$.
 - $\partial J / \partial x(t)$ can be gotten directly from the loss function.
 - $\partial x(t+1) / \partial x(t)$ only requires simulating one timestep forward, for each variable.

Since we have to simulate between each pair of timesteps t and $t + 1$, this is equivalent to running through the whole model once (per variable in x).

Once we've done that, we don't need to run the whole simulation for each $f(q)$: we only have to run one timestep, to see how it affects $x(q + 1)$.

We can think of this as "pre-simulating" the effect that our states have on the loss, so that we only have to see how $f(q)$ affects the first in that chain of timesteps: $x(q + 1)$.

10.8 Code

Our code is spread across 8 jupyter notebooks, including tutorials justifying each function, and demonstrating how to use it.

Here, we provide the file `helper.py`, which strips away these additional materials and simply provides the functions.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from PIL import Image
4 from matplotlib import gridspec
5 from scipy.interpolate import griddata
6 from matplotlib.colors import ListedColormap
7 from scipy import sparse
8 import copy
9
10 #NASA Logo
11
12 def png_to_matrix(file_path):
13     """
14     Convert a PNG image to a NumPy matrix.
15
16     Args:
17     file_path (str): Path to the PNG file.
18
19     Returns:
20     numpy.ndarray: 2D matrix representing the image.
21     """
22     # Open the image using PIL
23     img = Image.open(file_path)
24
25     # Convert the image to grayscale if it's not already
26     img = img.convert('L')
27
28     # Convert the image to a NumPy array
29     matrix = np.array(img)
30
31     return matrix
32
33 file_path = 'ImageProcessingDemo128.png'
34 NASA_map = png_to_matrix(file_path)[::-1,:]
35
36
37
38 ### Diffusion-Advection-Forcing Model ###
39
40 def make_M_2d_diffusion_advection_forcing(nr: int, nc: int, dt: float,
```

```

41         KX: np.ndarray, KY: np.ndarray,
42         DX_C: np.ndarray, DY_C: np.ndarray,
43         DX_G: np.ndarray, DY_G: np.ndarray,
44         VX: np.ndarray, VY: np.ndarray,
45         RAC: np.ndarray,
46         F : float,
47         cyclic_east_west: bool = True,
48         cyclic_north_south: bool = False,
49         M_is_sparse=True):
50
51     """
52     Creates linear model M which can be used to forward-simulate a
53     discrete approximation of a
54     2D diffusion-advection-forcing model.
55
56      $c(t+1) = Mc(t) + F f(t)$ 
57
58     Spatially-variant case
59
60     nr:      the number of rows of discrete cells.
61     nc:      the number of cols of discrete cells.
62     dt:      duration of a timestep
63
64     (i,j) = (0, 0) is the southwesternmost cell
65     (i,j) = (-1,-1) is the northeasternmost cell
66
67     For the below definitions:
68
69     KX:      the diffusivity constant matrix along the x-axis (between
70              columns)
71              KX[i,j] gives diffusivity at the boundary between cells
72                  [i,j-1] and [i,j]
73
74     KY:      the diffusivity constant matrix along the y-axis (between
75              rows)
76              KY[i,j] gives diffusivity at the boundary between cells
77                  [i-1,j] and [i,j]
78
79     DX_C:    the horizontal distance (x-axis) matrix between the centers
80              cells in adjacent columns
81              DX_C[i,j] gives the distance between the centers of cells
82                  [i,j-1] and [i,j]
83
84     DY_C:    the vertical distance (y-axis) matrix between the centers of
85              cells in adjacent rows
86              DY_C[i,j] gives the distance between the centers of cells
87                  [i-1,j] and [i,j]

```

```

79     DX_G:    the horizontal length (x-axis) matrix of a cell along one
            edge.
80             DX_G[i,j] gives the length of the "south" side of cell [i
                ,j]
81
82     DY_G:    the vertical length (y-axis) matrix of a cell along one edge.
83             DY_G[i,j] gives the length of the "west" side of cell [i,
                j]
84
85     VX:      the velocity constant matrix along the x-axis (between
            columns)
86             VX[i,j]   gives the velocity at the boundary between
                cells [i,j-1] and [i,j]
87
88     VY:      the velocity constant matrix along the y-axis (between rows)
89             VY[i,j]   gives the velocity at the boundary between
                cells [i-1,j] and [i,j]
90
91     RAC:     the area of a cell.
92             RAC[i,j]  gives the area of cell [i,j]
93
94     F:       the forcing term constant. This is the same for all cells.
95
96     cyclic_east_west:  if True, cell [i, 0] is east of cell [i,-1]
97
98     cyclic_north_south: if True, cell [0, j] is north of cell [-1,j]
99
100    M_is_sparse: if True, return a sparse matrix. If False, return a
            dense matrix.
101    """
102
103    if KX.shape != (nr, nc+1):
104        raise ValueError("KX doesn't have the right shape for your
            dimensions!")
105    if KY.shape != (nr+1, nc):
106        raise ValueError("KY doesn't have the right shape for your
            dimensions!")
107    if DX_C.shape != (nr, nc+1):
108        raise ValueError("DX_C doesn't have the right shape for your
            dimensions!")
109    if DY_C.shape != (nr+1, nc):
110        raise ValueError("DY_C doesn't have the right shape for your
            dimensions!")
111    if DX_G.shape != (nr+1, nc):
112        raise ValueError("DX_G doesn't have the right shape for your
            dimensions!")
113    if DY_G.shape != (nr, nc+1):

```

```

114         raise ValueError("DX_G doesn't have the right shape for your
115                             dimensions!")
116     if VX.shape != (nr, nc+1):
117         raise ValueError("VX doesn't have the right shape for your
118                             dimensions!")
119     if VY.shape != (nr+1, nc):
120         raise ValueError("VY doesn't have the right shape for your
121                             dimensions!")
122
123     size = nr * nc
124
125     if M_is_sparse:
126         M = sparse.lil_matrix((size, size))
127     else:
128         M = np.zeros((size, size))
129
130     beta = dt / RAC
131
132     # Shorthand variables
133
134     S = KX*DY_G/DX_C
135     T = KY*DX_G/DY_C
136
137     S_IJ, T_IJ = S[:, :-1], T[:-1, :]
138     S_IJP1, T_IP1J = S[:, 1:], T[1:, :]
139
140     R = VX*DY_G
141     Q = VY*DX_G
142
143     R_IJ, Q_IJ = R[:, :-1], Q[:-1, :]
144     R_IJP1, Q_IP1J = R[:, 1:], Q[1:, :]
145
146     # Contributions from diffusion (d)
147
148     d_IP1_J = beta * T_IP1J
149     d_IM1_J = beta * T_IJ
150     d_I_JP1 = beta * S_IJP1
151     d_I_JM1 = beta * S_IJ
152
153     d_IJ = - d_IP1_J - d_IM1_J - d_I_JP1 - d_I_JM1
154
155     # Contributions from advection (a)
156
157     a_IM1_J = beta * np.maximum(Q_IJ, 0)
158     a_IP1_J = - beta * np.minimum(Q_IP1J, 0)

```



```

158     a_I_JM1 = beta * np.maximum(R_IJ, 0)
159     a_I_JP1 = - beta * np.minimum(R_IJP1, 0)
160
161     a_IJ = beta * (\
162         np.minimum(Q_IJ, 0) - np.maximum(Q_IP1J, 0) + \
163         np.minimum(R_IJ, 0) - np.maximum(R_IJP1, 0))
164
165
166     #Create array to store indices
167
168     c = np.zeros([nr,nc])
169     c_indices = np.arange(len(c.ravel()))
170     c_indices = np.array(np.reshape(c_indices, [nr, nc]))
171
172     for i in range(nr): #y-axis (north, south)
173         for j in range(nc): #x-axis (east, west)
174
175             #Get current position
176             ind_here = c_indices[i,j]
177
178             # Currently we have no adjacent cells, we need to populate
179             # them
180             ind_N = np.nan
181             ind_E = np.nan
182             ind_S = np.nan
183             ind_W = np.nan
184
185             # Get indices for each direction
186             # south
187             if i > 0:
188                 ind_S = c_indices[i-1, j]
189             elif cyclic_north_south:
190                 ind_S = c_indices[-1, j]
191
192             # north
193             if i < nr-1:
194                 ind_N = c_indices[i+1, j]
195             elif cyclic_north_south:
196                 ind_N = c_indices[0, j]
197
198             # west
199             if j > 0:
200                 ind_W = c_indices[i, j-1]
201             elif cyclic_east_west:
202                 ind_W = c_indices[i, -1]
203
204             # east

```

```

204         if j < nc-1:
205             ind_E = c_indices[i, j+1]
206         elif cyclic_east_west:
207             ind_E = c_indices[i, 0]
208
209         # Now that we have our indices, we can fill in our matrix
210
211         M[ind_here, ind_here] = 1 + d_IJ[i,j] + a_IJ[i,j] \
212             - F #Forcing term is the same for all
213                 cells
214
215         if np.isfinite(ind_W):
216             # cell to the west
217             M[ind_here, ind_W] = 0 + d_I_JM1[i,j] + a_I_JM1[i,j]
218         if np.isfinite(ind_E):
219             # cell to the east
220             M[ind_here, ind_E] = 0 + d_I_JP1[i,j] + a_I_JP1[i,j]
221         if np.isfinite(ind_N):
222             # cell to the north
223             M[ind_here, ind_N] = 0 + d_IP1_J[i,j] + a_IP1_J[i,j]
224         if np.isfinite(ind_S):
225             # cell to the south
226             M[ind_here, ind_S] = 0 + d_IM1_J[i,j] + a_IM1_J[i,j]
227
228         if sparse:
229             M = M.tocsr()
230
231         return M
232
233
234
235
236 ### Simulate Model ###
237
238 def compute_linear_time_evolution(c0, M, saved_timesteps, duration,
239     debug = False):
240     """
241     Compute linear time evolution of model:  $c(t+1) = Mc(t)$ 
242
243     Args:
244     c0 (array): Initial state vector (a,1)
245     M (array): Linear model matrix (a,a)
246     saved_timesteps (list): Timesteps to save state
247     duration (int): Number of timesteps to simulate
248     debug (bool): If True, print progress every 10 steps
249

```

```

250     Returns:
251     saved_timesteps (list): List of timesteps where state was saved
252     saved (list): List of state vectors at each saved timestep
253     """
254     r = np.zeros_like(c0)
255     return compute_affine_time_evolution(c0, M, r, saved_timesteps,
256                                         duration, debug)
257
258 def compute_linear_time_evolution_simple(c0, M, num_saved_timesteps,
259                                         duration_per_saved_timestep,
260                                         debug = False):
261     """
262     Compute linear time evolution of model:  $c(t+1) = Mc(t)$ 
263
264     Simplified with evenly spaced saved_timesteps.
265
266     Args:
267     c0 (array): Initial state vector (a,1)
268     M (array): Linear model matrix (a,a)
269     num_saved_timesteps (int): Number of saved_timesteps to save
270     duration_per_saved_timestep (int): Number of timesteps between
271         saved_timesteps
272     debug (bool): If True, print progress every 10 steps
273
274     Returns:
275     list: 1D arrays representing saved_timesteps in time
276     """
277     saved_timesteps = [i*duration_per_saved_timestep for i in range(
278         num_saved_timesteps)] # Evenly spaced saved_timesteps
279     duration = num_saved_timesteps * duration_per_saved_timestep + 1 #
280         Duration must be longer than the last saved timestep
281
282     return compute_linear_time_evolution(c0, M, saved_timesteps, duration
283                                         , debug)
284
285 def compute_affine_time_evolution(c0, M, r, saved_timesteps, duration,
286                                   debug = False):
287     """
288     Compute affine time evolution model with affine term:  $c(t+1) = Mc(t)$ 
289         + r
290
291     Args:
292     c0 (array): Initial state vector (a,1)
293     M (array): Linear model matrix (a,a)
294     r (array): Affine term vector (a,1)
295     saved_timesteps (list): Timesteps to save state
296     duration (int): Number of timesteps to simulate

```



```
331
332
333
334 ### Plotting Heatmaps ###
335
336 def plot_multi_heatmap_time_evolution(saved_timesteps,
337                                     many_states_over_time,
338                                     nr, nc, titles, big_title,
339                                     vmin=None, vmax=None, is_ld=False):
340     """
341     Display time evolution of multiple 1D or 2D arrays over time, side by
342     side.
343
344     Args:
345     saved_timesteps (list or None): List of times that we save our state,
346     or None if not shown
347     many_states_over_time (list of lists): Each inner list contains 1D
348     arrays of state at each timestep
349     nr (int): Number of rows in each 2D array (or 1 for 1D state)
350     nc (int): Number of columns in each 2D array (or length of 1D array)
351     titles (list): Titles for each subplot
352     big_title (str): Overall title for the entire plot
353     vmin, vmax (float, optional): Min/max values for color scaling
354     is_ld (bool): Whether the input state is 1D (True) or 2D (False)
355
356     Displays heatmaps of several 1D or 2D arrays evolving over time.
357     """
358     # Reconstruct 1D arrays into 2D arrays
359     if is_ld:
360         plottable_states = [[state.reshape(1, -1) for state in
361                             state_over_time]
362                             for state_over_time in many_states_over_time]
363     else:
364         plottable_states = [[state.reshape(nr, nc) for state in
365                             state_over_time]
366                             for state_over_time in many_states_over_time]
367
368     # Calculate global min and max for consistent color scaling if not
369     provided
370     if vmin is None:
371         vmin = min(np.min(state) for state_over_time in
372                   many_states_over_time for state in state_over_time)
373     if vmax is None:
374         vmax = max(np.max(state) for state_over_time in
375                   many_states_over_time for state in state_over_time)
376
377     # Create the heatmaps
```

```

369 num_states = len(many_states_over_time)
370 num_timesteps = len(many_states_over_time[0]) # Assume all state
      runs have the same number of timesteps
371
372 # Calculate figure size and height ratios
373 title_height = 0.5 # inches
374 subplot_height = 2 if is_1d else 5 # inches
375 total_height = title_height + (subplot_height * num_timesteps)
376 fig_width = 5 * num_states + 1
377
378 # Create figure with two subfigures
379 fig = plt.figure(figsize=(fig_width, total_height))
380 subfigs = fig.subfigures(2, 1, height_ratios=[title_height,
      subplot_height * num_timesteps])
381
382 # Add the main title to the top subfigure
383 subfigs[0].suptitle(big_title, fontsize=16)
384
385 # Create gridspec for the bottom subfigure (plot grid)
386 gs = subfigs[1].add_gridspec(num_timesteps, num_states + 1,
      width_ratios=[1]*num_states + [0.05])
387
388 for i, state_over_time in enumerate(plottable_states):
389     for j, state in enumerate(state_over_time):
390         ax = subfigs[1].add_subplot(gs[j, i])
391         im = ax.imshow(state, aspect='auto', cmap='coolwarm', vmin=
            vmin, vmax=vmax, origin='lower')
392
393         # Set the title, including timestep if provided
394         if saved_timesteps is None:
395             ax.set_title(f'{titles[i]}')
396         else:
397             ax.set_title(f'{titles[i]}\nt={saved_timesteps[j]}')
398
399         ax.set_xlabel('$x$')
400         if not is_1d:
401             ax.set_ylabel('$y$')
402         else:
403             ax.set_yticks([])
404
405         # Add colorbar for each row
406         if i == num_states - 1: # Only for the last column
407             cbar_ax = subfigs[1].add_subplot(gs[j, -1])
408             plt.colorbar(im, cax=cbar_ax)
409
410 plt.tight_layout()
411 plt.show()

```

```
412
413 def plot_1d_heatmap_time_evolution(saved_timesteps, state_over_time):
414     """
415     Displays the time evolution of a 1D state as a series of heatmaps.
416
417     Args:
418     saved_timesteps (list): List of time values corresponding to each
419         state array
419     state_over_time (list): List of 1D numpy arrays, each representing
420         the state at a time saved
421
422     Returns:
423     None: Displays the plot using matplotlib
424
425     """
426     plot_multi_heatmap_time_evolution(
427         saved_timesteps=saved_timesteps,
428         many_states_over_time=[state_over_time],
429         nr=1,
430         nc=len(state_over_time[0]),
431         titles=["Evolution in time"],
432         big_title="Evolution in time"
433     )
434
435
436 def plot_2d_heatmap_time_evolution(saved_timesteps, state_over_time, nr,
437     nc, vmin = None, vmax = None):
438     """
439     Displays the time evolution of 2D state as a series of heatmaps.
440
441     Args:
442     saved_timesteps (list): List of time values corresponding to each
443         state array in time
444     state_over_time (list): List of 1D numpy arrays, each representing
445         the state at a time saved
446     nr (int): Number of rows in the 2D grid
447     nc (int): Number of columns in the 2D grid
448     vmin (float, optional): Minimum value for color scaling. If None,
449         calculated from state.
450     vmax (float, optional): Maximum value for color scaling. If None,
451         calculated from state.
452
453     Returns:
454     None: Displays the plot using matplotlib
455
456     """
457     # We're passing a single state over time, so we wrap it in another
```

```

    list
452 plot_multi_heatmap_time_evolution(
453     saved_timesteps=saved_timesteps,
454     many_states_over_time=[state_over_time],
455     nr=nr,
456     nc=nc,
457     titles=["Evolution in time"],
458     big_title="Evolution in time",
459     vmin=vmin,
460     vmax=vmax
461 )
462
463
464 def plot_multi_heatmap(many_states, nr, nc, titles, big_title, vmin=None,
465     vmax=None):
466     """
467     Display heatmaps of multiple 2D states side by side for comparison.
468
469     Args:
470     many_states (list): List of 1D arrays to be reshaped into 2D
471     nr, nc (int): Number of rows and columns for reshaping
472     titles (list): Titles for each heatmap
473     big_title (str): Overall title for the plot
474     vmin, vmax (float, optional): Min/max values for color scaling
475
476     Displays heatmaps side by side for comparison
477     """
478     # We're comparing at a single timestep, so we wrap each data array in
479     # its own list
480     many_states_over_time = [[state] for state in many_states]
481
482     plot_multi_heatmap_time_evolution(
483         saved_timesteps=None, # Single timestep
484         many_states_over_time=many_states_over_time,
485         nr=nr,
486         nc=nc,
487         titles=titles,
488         big_title=big_title,
489         vmin=vmin,
490         vmax=vmax
491     )
492
493 ### Covariance functions ###
494
495 def compute_covariance_gaussian_dropoff(a, b, std_dev = 1):

```



```

496     """
497     Compute Gaussian covariance matrix with exponential dropoff.
498
499     Calculates covariance using  $e^{\{-d / 2s\}}$ , where d is squared
500     Euclidean distance between a and b, and s is the standard deviation
501     std_dev.
502
503     Args:
504     a, b (array-like): Input vectors, represents spatial coordinates
505     std_dev (float): Standard deviation of gaussian dropoff (default=1)
506
507     Returns:
508     numpy.ndarray: Covariance matrix
509     """
510
511     d = np.linalg.norm(a - b, axis=-1)**2          #Distance between a and b
512     covariance = np.exp( -d / (2*std_dev))          #Covariance matrix
513
514     return covariance
515
516 def vector_of_2d_indices(nr, nc):
517     """
518     Convert 2D array of indices to a 1D vector representation containing
519     the same indices.
520
521     Args:
522     nr (int): Number of rows
523     nc (int): Number of columns
524
525     Returns:
526     numpy.ndarray: Shape (nr*nc, 2), each row is [row, col] index
527     """
528     y, x = np.mgrid[:nr, :nc] #Get meshgrid of all of our indices
529     vector_form = np.column_stack((y.ravel(), x.ravel())) #Stack
530     return vector_form
531
532 def compute_covariance_matrix_gaussian_dropoff(nr, nc, std_dev=1):
533     """
534     Compute covariance matrix with Gaussian dropoff for 2D grid.
535
536     Compares all pairs of 2D indices and computes covariance using
537     exponential dropoff function  $e^{\{-d / 2s\}}$ , where
538     - d is squared Euclidean distance between a and b
539     - s is the standard deviation std_dev.
540

```

```

541     Args:
542     nr, nc (int): Number of rows and columns in 2D grid
543     std_dev (float): Standard deviation for Gaussian dropoff (default=1)
544
545     Returns:
546     np.ndarray: Covariance matrix of size (nr*nc, nr*nc)
547     """
548     # Get all 2D indices as a 1D vector of (row, col) pairs
549     index_vector = vector_of_2d_indices(nr, nc)
550
551     # Prepare indices for broadcasting
552     indices_i = index_vector[:, np.newaxis, :] # Shape: (nr*nc, 1, 2)
553     indices_j = index_vector[np.newaxis, :, :] # Shape: (1, nr*nc, 2)
554
555     # Compute covariance matrix
556     covariance_matrix = compute_covariance_gaussian_dropoff(
557         indices_i,
558         indices_j,
559         std_dev=std_dev
560     )
561
562     return covariance_matrix
563
564
565
566 ### Generating Smooth Data ###
567
568 def add_random_circles(matrix, num_circles, radius, values):
569     """
570     Add random circles to a matrix for interesting initial conditions.
571
572     Args:
573     matrix (numpy.ndarray): 2D array to modify
574     num_circles (int): Number of circles to add
575     radius (int): Radius of circles
576     values (list): Possible values for circles
577
578     Returns:
579     numpy.ndarray: Modified matrix with added circles
580     """
581
582     nr, nc = matrix.shape
583     for _ in range(num_circles):
584         center_r = np.random.randint(0, nr)
585         center_c = np.random.randint(0, nc)
586         r, c = np.ogrid[:nr, :nc]
587         mask = ((r - center_r)**2 + (c - center_c)**2 <= radius**2)

```

```

588         value = np.random.choice(values)
589         matrix[mask] = value
590     return matrix
591
592
593 def generate_random_vectors_mean_0_cov_C(nr, nc, C, num_vectors):
594     """
595     Generates random vectors from a normal distribution with mean 0 and
596     covariance C.
597
598     Args:
599     nr, nc (int): Dimensions of the 2D grid
600     C (np.ndarray): Covariance matrix
601     num_vectors (int): Number of random vectors to generate
602
603     Returns:
604     tuple: (zs, Lzs)
605         zs (list): Random normal vectors with mean 0 and covariance I
606              $z \sim N(0, I)$ 
607         Lzs (list): Random vectors with mean 0 and covariance C
608              $Lz \sim N(0, C)$ 
609     """
610     size = nr * nc #Size of the grid
611     zs = [np.random.randn(size, 1) for _ in range(num_vectors)] #Random
612         normal vectors with mean 0 and covariance I
613
614     # If  $z \sim N(0, I)$ , then  $Lz \sim N(0, C)$ 
615     L = np.linalg.cholesky(C)
616     Lzs = [L @ z for z in zs] #Random vectors with mean 0 and covariance
617         C
618
619     return zs, Lzs
620
621
622 def generate_random_vector_mean_0_cov_C(nr, nc, C):
623     """
624     Generates a random vector from a normal distribution with mean 0 and
625     covariance C.
626
627     Args:
628     nr, nc (int): Dimensions of the 2D grid
629     C (np.ndarray): Covariance matrix
630
631     Returns:
632     tuple: (z, Lz)
633         z (np.ndarray): Random normal vector with mean 0 and covariance I
634              $z \sim N(0, I)$ 
635         Lz (np.ndarray): Random vector with mean 0 and covariance C

```

```

631         Lz ~ N(0, C)
632     """
633     zs, Lzs = generate_random_vectors_mean_0_cov_C(nr, nc, C, 1)
634     return zs[0], Lzs[0]
635
636
637
638
639
640 def generate_true_and_first_guess_field_uniform_cov(C, nr, nc, gamma):
641     """
642     Generates two fields: a true field and a first-guess field, both with
643         the same covariance C.
644
645     Args:
646     C (np.ndarray): Covariance matrix for both fields
647     nr (int): Number of rows in the grid
648     nc (int): Number of columns in the grid
649     gamma (float): Fraction of the field shared between true and first-
650         guess fields
651
652     Returns:
653     tuple: (true_field, first_guess_field)
654         true_field (np.ndarray): Random field with covariance C
655         first_guess_field (np.ndarray): Random field sharing a component
656             with true_field, covariance C
657
658     Notes:
659     - f2 replaces f1 in the "first-guess" field to represent inaccuracies
660     - The shared component (f0) represents the fraction of the field that
661         is
662         common between the true and first-guess fields
663     """
664     f0, f1, f2 = generate_random_vectors_mean_0_cov_C(nr, nc, C, 3)[1]
665
666     true_field = f0 * gamma + f1 * (1-gamma)
667     first_guess_field = f0 * gamma + f2 * (1-gamma)
668
669     return true_field, first_guess_field
670
671 def generate_true_and_first_guess_field(C_known, C_error, nr, nc):
672     """
673     Generates two fields: a true field and a first-guess field, both with
674         specified covariances.
675
676     Args:
677     C_known (np.ndarray): Covariance matrix for the known part of the

```

```

        field
673 C_error (np.ndarray): Covariance matrix for the error part of the
        field
674 nr, nc (int): Dimensions of the 2D grid
675
676 Returns:
677 tuple: (true_field, first_guess_field)
678     true_field (np.ndarray): Random field with covariance C_known +
        C_error
679     first_guess_field (np.ndarray): Random field sharing a component
        with true_field, covariance C_known + C_error
680
681 Notes:
682 - f2 replaces f1 in the "first-guess" field to represent inaccuracies
683 - The shared component (f0) represents the fraction of the field that
        is
684     common between the true and first-guess fields
685 """
686 f0 = generate_random_vectors_mean_0_cov_C(nr, nc, C_known, 1)[1][0]
687 f1, f2 = generate_random_vectors_mean_0_cov_C(nr, nc, C_error, 2)[1]
688
689 true_field = f0 + f1
690 first_guess_field = f0 + f2
691
692 return true_field, first_guess_field
693
694
695 def generate_gaussian_field(n, nrv, ncv):
696     """
697     Randomly generates a 2D field composed of n Gaussian functions with
        distinct means and standard deviations.
698
699     Args:
700     n (int): Number of Gaussian functions to generate
701     nrv (int): Number of rows in the field
702     ncv (int): Number of columns in the field
703
704     Returns:
705     np.ndarray: A 2D array representing the generated Gaussian field
706
707     Notes:
708     - The field is made pseudo-periodic by creating three copies of each
        Gaussian function
709     """
710     mux = np.random.choice(ncv, n)
711     muy = np.random.choice(range(2, nrv - 2), n)
712     sigma_max = np.random.uniform(1, ncv/4, n)

```

```

713     sigmay = np.random.uniform(1,nrv/4,n)
714
715     #Combine all the gaussian functions to get the field
716
717     v = np.zeros((nr,nr))
718     for i in range(n):
719         for x in range(nrv):
720             for y in range(nrv):
721                 #We create three copies of our gaussian so that we get a
722                 #pseudo-periodic field
723
724                 # Original Gaussian
725                 gauss = np.exp(-((x-mux[i])**2/(2*sigmax[i]**2) + (y-muy
726                 [i])**2/(2*sigmay[i]**2)))
727
728                 # Shifted left
729                 gauss += np.exp(-((x-(mux[i]-ncv))**2/(2*sigmax[i]**2) +
730                 (y-muy[i])**2/(2*sigmay[i]**2)))
731
732                 # Shifted right
733                 gauss += np.exp(-((x-(mux[i]+ncv))**2/(2*sigmax[i]**2) +
734                 (y-muy[i])**2/(2*sigmay[i]**2)))
735
736                 v[y,x] += gauss
737
738     return v
739
740 def generate_circular_field(v):
741     """
742     Generates a circular field by taking the gradient of the input field
743     and rotating it by 90 degrees.
744
745     Args:
746     v (np.ndarray): Input 2D field
747
748     Returns:
749     tuple: (grad_v_x, grad_v_y)
750             grad_v_x (np.ndarray): X-component of the circular field
751             grad_v_y (np.ndarray): Y-component of the circular field
752     """
753     grad_v_y, grad_v_x = np.gradient(v)
754
755     return -grad_v_y, grad_v_x
756
757 def create_random_model(nr, nc, dt, F,
758                         num_gauss = 16,
759                         DX_C = None, DY_C = None, DX_G = None, DY_G =

```

```

None, RAC = None,
cyclic_east_west=True, cyclic_north_south=False):
"""
Creates a random model with a new velocity field and diffusivity
field.

- Velocity field is generated from a field of circular patterns, in
order to create a field
with low divergence.

- Diffusivity field is generated randomly, with 0 diffusivity on the
boundaries.

Args:
nr (int): Number of rows in the grid
nc (int): Number of columns in the grid
dt (float): Time step
F (float): Forcing coefficient
num_gauss (int, optional): Number of Gaussian functions for velocity
field generation. Defaults to 16.
DX_C, DY_C, DX_G, DY_G, RAC (np.ndarray, optional): Grid spacing and
area parameters. If None, set to arrays of ones.
cyclic_east_west (bool, optional): If True, applies cyclic conditions
east-west. Defaults to True.
cyclic_north_south (bool, optional): If True, applies cyclic
conditions north-south. Defaults to False.

Returns:
tuple: (M, params)
M (np.ndarray): Model matrix for 2D diffusion-advection-forcing
params (dict): Dictionary of parameters used to create the model

"""
# If none, just set everything to appropriately-sized array of 1's
if DX_C is None:
    DX_C = np.ones((nr, nc+1))
if DY_C is None:
    DY_C = np.ones((nr+1, nc))
if DX_G is None:
    DX_G = np.ones((nr+1, nc))
if DY_G is None:
    DY_G = np.ones((nr, nc+1))
if RAC is None:
    RAC = np.ones((nr, nc))

```

```

794     # Randomly generate diffusivities: must be positive
795     KX = np.random.rand(nr, nc+1)
796     KY = np.random.rand(nr+1, nc)
797     KX = np.abs(KX)
798     KY = np.abs(KY)
799
800     # Randomly generate velocities as above
801     gauss = generate_gaussian_field(num_gauss, nr+1, nc+1)
802     VX, VY = generate_circular_field(gauss)
803
804     # Create the model matrix
805     params = {
806         'nr': nr,
807         'nc': nc,
808         'dt': dt,
809         'KX': KX,
810         'KY': KY,
811         'DX_C': DX_C,
812         'DY_C': DY_C,
813         'DX_G': DX_G,
814         'DY_G': DY_G,
815         'VX': 100*VX[:-1, :],
816         'VY': 100*VY[:, :-1],
817         'RAC': RAC,
818         'F': F,
819         'cyclic_east_west': cyclic_east_west,
820         'cyclic_north_south': cyclic_north_south
821     }
822
823     M = make_M_2d_diffusion_advection_forcing(**params)
824
825     return M, params
826
827 def create_random_initial_ocean_state(nr, nc, C, num_circles, radius,
828 values):
829     """
830     Creates a random initial ocean state with specified covariance matrix
831     .
832
833     Args:
834     nr (int): Number of rows in the grid
835     nc (int): Number of columns in the grid
836     C (np.ndarray): Covariance matrix for the initial state
837     num_circles (int): Number of random circles to add
838     radius (int): Radius of circles
839     values (list): Possible values for circles

```



```

839 Returns:
840 tuple: (z, Lz)
841     z (np.ndarray): Random initial state with covariance C
842     Lz (np.ndarray): Random initial state with covariance
843     """
844
845     z = np.random.rand(nr,nc)
846     z = add_random_circles(z, num_circles, radius, values)
847     z = z.reshape((nr*nc,1))
848
849     L = np.linalg.cholesky(C)
850     Lz = L @ z
851
852     return z, Lz
853
854 def generate_world(nr, nc, dt, F, num_gauss=16, num_circles=20, radius=5,
855     values=[2,-2], std_dev=2):
856     """
857     Generates a world with an ocean state, atmosphere, and model matrix.
858
859     Args:
860     nr (int): Number of rows in the grid
861     nc (int): Number of columns in the grid
862     dt (float): Time step
863     F (float): Forcing parameter
864     num_gauss (int, optional): Number of Gaussian functions for velocity
865         field generation. Defaults to 16.
866     num_circles (int, optional): Number of circles for ocean state
867         generation. Defaults to 20.
868     radius (int, optional): Radius of circles for ocean state generation.
869         Defaults to 5.
870     values (list, optional): Values of circles for ocean state generation
871         . Defaults to [2,-2].
872     std_dev (int, optional): Standard deviation for Gaussian dropoff.
873         Defaults to 2.
874
875     Returns:
876     tuple: (C, c0, f, M)
877         C (np.ndarray): Covariance matrix
878         c0 (np.ndarray): Initial ocean state
879         f (np.ndarray): Atmosphere
880         M (np.ndarray): Model matrix
881     """
882     # Generate covariance matrix
883     C = compute_covariance_matrix_gaussian_dropoff(nr, nc, std_dev)
884
885     # Generate model matrix

```

```
880     M, params = create_random_model(nr, nc, dt, F, num_gauss=num_gauss)
881
882     # Generate initial ocean state
883     _, c0 = create_random_initial_ocean_state(nr, nc, C, num_circles=
        num_circles, radius=radius, values=values)
884
885     # Generate atmosphere
886     _, f = generate_random_vector_mean_0_cov_C(nr, nc, C)
887
888     return C, c0, f, M
889
890 #C, c0, f, M = generate_world(50, 50, 0.1, 1, num_gauss=16, num_circles
    =20, radius=5, values=[2,-2], std_dev=2)
891
892 # Get magnitude of f
893 #f = f/3
894 #f_mag = np.linalg.norm(f)
895 #print(f'Magnitude of f: {f_mag:.2f}')
896
897
898 ### Observe field ###
899
900 def observe(real, sigma, num_observations):
901     """
902     Generates noisy observations of a true state at randomly selected
        indices.
903
904     Args:
905     real (np.ndarray): The true state of the system, as a 1D array
906     sigma (float): The standard deviation of the observation noise
907     num_observations (int): The number of observations to make
908
909     Returns:
910     tuple: (indices, observations)
911         indices (np.ndarray): Array of randomly selected indices for
            observation
912         observations (np.ndarray): Noisy observations of the true state
            at the selected indices
913
914     Notes:
915     - Observations are made by adding Gaussian noise to the true state
        values
916     - The noise is generated as a 2D column vector
917     - Indices are selected without replacement, ensuring unique
        observation points
918     """
919     # We randomly select which indices to observe
```

```
920     indices = np.random.choice(len(real), num_observations, replace=False
921                                 )
922
923     # We observe the true state plus Gaussian noise
924     noise = np.random.normal(0, sigma, (num_observations, 1))
925
926     observations = real[indices] + noise
927
928     return indices, observations
929
930 def fill_nan_map_with_observations(indices, observations, nr, nc):
931     """
932     Maps observations to their corresponding positions in a 2D grid,
933     filling the rest with NaNs.
934
935     Args:
936     indices (np.ndarray): Indices of the observations in the flattened
937                           grid
938     observations (np.ndarray): Observed values
939     nr (int): Number of rows in the grid
940     nc (int): Number of columns in the grid
941
942     Returns:
943     np.ndarray: 2D array with observations at their corresponding
944                 positions and NaNs elsewhere
945     """
946     observed_state_2d = np.full((nr, nc), np.nan)
947     observed_state_2d.flat[indices] = observations.flatten()
948     return observed_state_2d
949
950 def interpolate_observation_map(observed_state_2d, extend=False):
951     """
952     Interpolates a 2D grid of observed values, to fill in NaN values (
953         representing unobserved points).
954
955     Args:
956     observed_state_2d (np.ndarray): 2D array with observed values and
957                                     NaNs
958     extend (bool, optional): If True, extends the observed state by three
959                             copies horizontally. Defaults to False.
960
961     Returns:
962     np.ndarray: 2D array of interpolated values
963
964     Notes:
965     - If extend is True:
966       - The observed state is extended by three copies horizontally
```

```
960     - Interpolation is performed on the extended grid
961     - The middle third of the interpolated result is returned
962 - If extend is False:
963     - Interpolation is performed on the original grid
964 - Linear interpolation is used, with NaN values for points outside
    the convex hull of observations
965 """
966 nr, nc = observed_state_2d.shape
967
968 if extend:
969     # Extend state by three copies
970     observed_state_2d = np.hstack((observed_state_2d,
971                                     observed_state_2d, observed_state_2d))
972
973 extended_nr, extended_nc = observed_state_2d.shape
974
975 # Create grid coordinates
976 x, y = np.meshgrid(np.arange(extended_nc), np.arange(extended_nr))
977
978 # Find non-NaN indices and values
979 observed_indices = np.where(~np.isnan(observed_state_2d.flatten()))
980 [0]
981 observed_values = observed_state_2d.flatten()[observed_indices]
982
983 # Create points and grid for interpolation
984 points = np.column_stack((x.flat[observed_indices], y.flat[
985     observed_indices]))
986 grid_x, grid_y = np.meshgrid(np.arange(extended_nc), np.arange(
987     extended_nr))
988
989 # Perform interpolation
990 interpolated_2d = griddata(points, observed_values, (grid_x, grid_y),
991     method='linear', fill_value=np.nan)
992
993 if extend:
994     # Extract the middle third
995     return interpolated_2d[:, nc:2*nc]
996 else:
997     return interpolated_2d
998
999 def observe_over_time(ocean_states, sigma, num_obs_per_timestep, nr, nc):
1000     """
1001     Observes the ocean state at each timestep, and places those
1002     observations in
1003     a 2d array with NaNs for unobserved points.
```

```

1000     Args:
1001     ocean_states (list): List of ocean states at each timestep
1002     sigma (float): Standard deviation of observation noise
1003     num_obs_per_timestep (int): Number of observations per timestep
1004     nr (int): Number of rows in the grid
1005     nc (int): Number of columns in the grid
1006
1007     Returns:
1008     list: List of observed ocean states at each timestep, each as a 2D
1009           array with NaNs for unobserved points
1010
1011     """
1012     # Observe ocean state at each timestep
1013     indices_and_observations_over_time = [observe(ocean_state, sigma,
1014                                                  num_obs_per_timestep)
1015                                           for ocean_state in
1016                                           ocean_states]
1017
1018     # Place on a map: unobserved points are filled with NaN
1019     observed_state_over_time_2d = [fill_nan_map_with_observations(indices
1020                                                                    , observations_t, nr, nc)
1021                                    for indices, observations_t in
1022                                    indices_and_observations_over_time
1023                                    ]
1024
1025     return observed_state_over_time_2d
1026
1027 ##### Compute Adjoints #####
1028
1029 def compute_Jt(xt_true, xt_guess):
1030     """
1031     Computes squared loss between two vectors at time t.
1032
1033     Args:
1034     xt_true (np.ndarray): True state vector at time t
1035     xt_guess (np.ndarray): Guessed state vector at time t
1036
1037     Returns:
1038     float: Squared loss, or 0 if no valid terms
1039     """
1040     # Sum over all valid terms, using numpy to treat nans as zeros
1041     result = np.nansum((xt_true - xt_guess)**2)
1042     if np.isnan(result):

```

```
1041         return 0
1042     else:
1043         return result
1044
1045 def compute_J(x_true, x_guess):
1046     """
1047     Computes total squared loss between two vectors across all timesteps.
1048
1049     Args:
1050     x_true (list): List of true state vectors at each timestep
1051     x_guess (list): List of guessed state vectors at each timestep
1052
1053     Returns:
1054     float: Total squared loss across all timesteps
1055     """
1056     return np.sum([
1057         compute_Jt(x_true[i], x_guess[i]) for i in range(len(
1058             x_true))
1059     ])
1060
1061 def compute_DJ_Dxt(xt_true, xt_guess):
1062     """
1063     Computes partial derivative of squared loss w.r.t. guessed state at
1064     time t.
1065
1066     Args:
1067     xt_true (np.ndarray): True state vector at time t
1068     xt_guess (np.ndarray): Guessed state vector at time t
1069
1070     Returns:
1071     np.ndarray: Partial derivative of loss, with NaNs treated as 0
1072     """
1073     return np.nan_to_num( 2*(xt_guess - xt_true), nan = 0 )
1074
1075 def compute_adjoints(DJ_Dx, dxtp1_dxt):
1076     """
1077     Computes adjoint variables for optimization using backwards-time
1078     recursion.
1079
1080     Args:
1081     DJ_Dx (list): List of partial derivatives of loss w.r.t. state at
1082         each timestep
1083     dxtp1_dxt (list): List of total derivatives of next state w.r.t.
1084         current state at each timestep
1085
1086     Returns:
```

```

1083     list: Adjoint variables for each timestep, in forward time order
1084     """
1085
1086     tau = len(DJ_Dx)
1087     adjoints = [0] * tau # Initialize list of adjoints
1088
1089     adjoints[tau-1] = DJ_Dx[tau-1]
1090
1091     for t in range(tau-2, -1, -1): # Backwards in time
1092         adjoint = DJ_Dx[t] + dxtp1_dxt[t].dot( adjoints[t+1] )
1093
1094         adjoints[t] = adjoint
1095
1096     return adjoints
1097
1098 def compute_dJ_df(M, F, observed_state_over_time,
1099                  simulated_state_over_time):
1100     """
1101     Computes the gradient of the loss with respect to the forcing field f
1102         for the linear model:
1103         x(t+1) = Mx(t) + Ff
1104
1105     Args:
1106     M (np.ndarray): Model matrix
1107     F (float): Forcing coefficient
1108     observed_state_over_time (list): List of observed states at each
1109         timestep
1110     simulated_state_over_time (list): List of simulated states at each
1111         timestep
1112
1113     Returns:
1114     np.ndarray: Gradient of the loss with respect to the forcing field f
1115     """
1116     num_timesteps = len(observed_state_over_time)
1117     vec_length = len(observed_state_over_time[0])
1118
1119     #Compute adjoints
1120     DJ_Dx = [compute_DJ_Dxt(observed_state_over_time[i],
1121                             simulated_state_over_time[i])
1122               for i in range(num_timesteps)] # partial J / partial x(t)
1123
1124     dxtp1_dxt = [M.T for i in range(num_timesteps-1)] #dx(t+1)/dx(t)
1125
1126     adjoints = compute_adjoints(DJ_Dx, dxtp1_dxt) # dJ/dx(t) = lambda(t)
1127
1128     # Compute gradient for each timestep: how f being applied at time t
1129     # affects J

```

```

1124     dJ_dft = [ F * adjoint for adjoint in adjoints[1:] ] # dJ/df(t) = dx(
1125         t+1)/df(t) dJ/dx(t+1)
1126
1127     #f is applied the same at all timesteps
1128     dJ_df = np.sum(dJ_dft, axis=0) # dJ/df = sum_t dJ/df(t)
1129     return dJ_df
1130
1131 ### Gradient Descent ###
1132
1133 losses_template = { #Losses at each iteration
1134     "ocean_misfit": [],
1135     "atmosphere_misfit": [],
1136     "mahalanobis(covariance similarity)": [],
1137 }
1138
1139
1140 def update_losses(losses, ocean_states_observed, ocean_states_simulated,
1141     f_guess, f_adjust, f_true, C_error):
1142     """
1143     Updates the losses dictionary with new loss values for ocean,
1144     atmosphere, and control adjustment.
1145
1146     Args:
1147     losses (dict): Dictionary containing lists of loss values
1148     ocean_states_observed (list): List of observed ocean states
1149     ocean_states_simulated (list): List of simulated ocean states
1150     f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1151     f_adjust (np.ndarray): Adjustment to the atmospheric forcing field
1152     f_true (np.ndarray): True atmospheric forcing field
1153     C_error (np.ndarray): Covariance matrix for the control error
1154
1155     Returns:
1156     dict: Updated losses dictionary with new loss values appended
1157     """
1158     f_i = f_guess + f_adjust
1159
1160     # Compute losses
1161     ocean_loss_i = compute_J(ocean_states_observed,
1162         ocean_states_simulated) # J_{ocean} = J
1163     atmos_loss_i = compute_Jt(f_true, f_i)
1164         # J_{atm} = misfit of atm
1165
1166     mahal = f_adjust.T @ np.linalg.inv(C_error) @ f_adjust # C_error
1167         should be the covariance of our adjustment
1168     mahal_loss_i = np.linalg.norm(mahal)

```



```

1165     #Store losses
1166     losses["ocean_misfit"].append(ocean_loss_i)
1167     losses["atmosphere_misfit"].append(atmos_loss_i)
1168     losses["mahalanobis(covariance similarity)"].append(mahal_loss_i)
1169
1170     return losses
1171
1172 possible_debug_vars = { #Debug variables to compute at each iteration
1173     "Norm of s_i": [],
1174     "Expected Delta J w simple gd": [],
1175     "Expected Delta J w update rule": [],
1176     "Norm of simple gd ui": [],
1177     "Norm of update rule ui": [],
1178     "Normalized dot product $a_i$ and $u_i$": [],
1179     "$s_i^T C s_i$": [],
1180     "$a_i$": [],
1181     "Normalized $a_i^T s_i$": [],
1182     "Normalized $-C s_i \cdot a_i$": [],
1183     "Norm of $a_i$": [],
1184     "Actual Delta J": []
1185 }
1186
1187 def update_debug_vars(debug_vars, x0, M, F, f_guess, f_adjust,
1188                      C_known, C_error,
1189                      s, step_size, ui,
1190                      ocean_states_simulated, ocean_states_observed):
1191     """
1192     Updates the debug variables dictionary with various metrics for
1193     gradient descent analysis.
1194
1195     Args:
1196     debug_vars (dict): Dictionary containing lists of debug variable
1197                        values
1198     x0 (np.ndarray): Initial ocean state
1199     M (np.ndarray): Model matrix
1200     F (float): Forcing coefficient
1201     f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1202     f_adjust (np.ndarray): Adjustment to the atmospheric forcing field
1203     C_known (np.ndarray): Covariance matrix for the known portion of the
1204                          control
1205     C_error (np.ndarray): Covariance matrix for the control error
1206     s (np.ndarray): Gradient of the loss with respect to the forcing
1207                     field
1208     step_size (float): Step size for gradient descent
1209     ui (np.ndarray): Update vector for the current iteration
1210     ocean_states_simulated (list): List of simulated ocean states
1211     ocean_states_observed (list): List of observed ocean states

```

```

1208
1209 Returns:
1210 dict: Updated debug_vars dictionary with new values appended to each
      metric
1211 """
1212 #Initialize useful variables
1213 num_timesteps = len(ocean_states_observed)
1214 ui_simple_gd = -step_size * s
1215 delta = s.T @ (ui_simple_gd)
1216
1217 #Compute debug vars
1218 norm_s = np.linalg.norm(s)
1219 exp_delta_J_simple_gd = (s.T @ ui_simple_gd)[0,0]
1220 exp_delta_J_update_rule = (s.T @ ui)[0,0]
1221 norm_simple_ui = np.linalg.norm(ui_simple_gd)
1222 norm_ui = np.linalg.norm(ui)
1223 norm_dot_product = (f_adjust.T @ ui)[0,0] / (np.linalg.norm(f_adjust)
      * np.linalg.norm(ui))
1224 sTCs = (s.T @ C_error @ s)[0,0]
1225 ai = f_adjust
1226 normalized_aiTs = (f_adjust.T @ s)[0,0] / (np.linalg.norm(f_adjust) *
      np.linalg.norm(s))
1227 normalized_Csdotai = (f_adjust.T @ (C_error @ s))[0,0] / (np.linalg.
      norm(f_adjust) * np.linalg.norm(C_error @ s))
1228 norm_ai = np.linalg.norm(f_adjust)
1229
1230 #Store debug vars
1231 debug_vars["Norm of s_i"].append(norm_s)
1232 debug_vars["Expected Delta J w simple gd"].append(
      exp_delta_J_simple_gd)
1233 debug_vars["Expected Delta J w update rule"].append(
      exp_delta_J_update_rule)
1234 debug_vars["Norm of simple gd ui"].append(norm_simple_ui)
1235 debug_vars["Norm of update rule ui"].append(norm_ui)
1236 debug_vars["Normalized dot product $a_i$ and $u_i$"].append(
      norm_dot_product)
1237 debug_vars["$s_i^T Cs_i$"].append(sTCs)
1238 debug_vars["$a_i$"].append(ai)
1239 debug_vars["Normalized $a_i^T s_i$"].append(normalized_aiTs)
1240 debug_vars["Normalized $-Cs_i \cdot a_i$"].append(-normalized_Csdotai
      )
1241 debug_vars["Norm of $a_i$"].append(norm_ai)
1242
1243 #Handle last debug var: Actual Delta J
1244 f_new = f_guess + f_adjust + ui
1245 _, new_ocean_states_simulated = compute_affine_time_evolution_simple(
      x0, M, F*f_new, num_timesteps)

```

```

1246
1247     new_J = compute_J(ocean_states_observed, new_ocean_states_simulated)
1248     old_J = compute_J(ocean_states_observed, ocean_states_simulated)
1249
1250     actual_delta_J = new_J - old_J
1251
1252     debug_vars["Actual Delta J"].append(actual_delta_J)
1253
1254     return debug_vars
1255
1256
1257 def gradient_descent_template(M, F, f_true, f_guess, C_known, C_error,
1258                             # World parameters
1259                             x0, num_timesteps,
1260                             #
1261                             # Simulation parameters
1262                             ocean_states_observed, num_iters, step_size
1263                             , # Optimization parameters
1264                             update_rule, update_params, disp=False):
1265     """
1266     Perform gradient descent to optimize the atmospheric forcing field.
1267
1268     The step we take at each iteration is computed using a modified
1269     update rule, and extra parameters as necessary.
1270
1271     Args:
1272     M: Model matrix
1273     F: Scalar constant for forcing
1274     f_true: True atmospheric forcing field
1275     f_guess: Initial guess for the atmospheric forcing field
1276     x0: Initial ocean state
1277     num_timesteps: Number of timesteps
1278     ocean_states_observed: Observed ocean states
1279     num_iters: Number of iterations
1280     step_size: Step size for gradient descent
1281     C_known: Covariance matrix for the known portion of the control
1282     C_error: Covariance matrix for the control error
1283     update_params: Function to compute the update rule to take at
1284                   each iteration
1285     extra_params: Extra parameters to pass to the update rule
1286     disp: Flag to print information
1287
1288     Returns:
1289     f: Optimized atmospheric forcing field
1290     losses: Dictionary of losses
1291           ocean_misfit: Ocean loss at each iteration
1292           atmosphere: Atmospheric loss at each iteration

```

```
1287         $a_iC^{-1}a_i$: Mahalanobis distance for the control
1288         adjustment at each iteration
1289
1288     """
1289     size = f_guess.shape[0]
1290     f_adjust = np.zeros((size,1))
1291
1292     losses = copy.deepcopy(losses_template)
1293     debug_vars = copy.deepcopy(possible_debug_vars)
1294
1295     for i in range(num_iters):
1296         if i%10==0 and disp:
1297             print("Iteration", i)
1298
1299         #Compute results of previous update rule
1300         f_i = f_guess + f_adjust #f_i = f_0 + a_i
1301         _, ocean_states_simulated = compute_affine_time_evolution_simple(
1302             x0, M, F*f_i, num_timesteps)
1303
1304         # Compute and store losses
1305         losses = update_losses(losses, ocean_states_observed,
1306                                ocean_states_simulated, f_guess, f_adjust, f_true, C_error)
1307
1308         # Compute and store debug variables
1309
1310         s = compute_dJ_df(M, F, ocean_states_observed,
1311                            ocean_states_simulated)
1312         ui = update_rule(i, s, step_size, f_adjust, *update_params) #
1313             Update rule
1314
1315         debug_vars = update_debug_vars(debug_vars, x0, M, F, f_guess,
1316                                        f_adjust,
1317                                        C_known, C_error,
1318                                        s, step_size, ui,
1319                                        ocean_states_simulated,
1320                                        ocean_states_observed)
1321
1322         # Apply update rule to f_adjust
1323
1324         f_adjust = f_adjust + ui
1325
1326     return f_adjust, losses, debug_vars
1327
1328 def simple_gradient_update_rule(curr_iter, s, step_size, f_adjust):
```

```

1327     """
1328     Computes the update step for simple gradient descent.
1329
1330     Args:
1331     curr_iter (int): Current iteration number (unused in this function)
1332     s (np.ndarray): Gradient of the loss with respect to the forcing
        field
1333     step_size (float): Step size for gradient descent
1334     f_adjust (np.ndarray): Current adjustment to the forcing field (
        unused in this function)
1335
1336     Returns:
1337     np.ndarray: Update step for the forcing field adjustment
1338     """
1339     return -step_size * s # Just use the gradient of the loss
1340
1341 def simple_gradient_descent(M, F, f_true, f_guess, C_known, C_error,
        # World parameters
1342                             x0, timesteps,
        #
        # Simulation parameters
1343                             ocean_states_observed, num_iters, step_size,
        # Optimization parameters
1344                             disp=False):
        #
        # Optimization method
1345     """
1346     Implements simple gradient descent for optimizing the atmospheric
        forcing field.
1347
1348     Args:
1349     M (np.ndarray): Model matrix
1350     F (float): Forcing coefficient
1351     f_true (np.ndarray): True atmospheric forcing field
1352     f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1353     C_known (np.ndarray): Covariance matrix for the known portion of the
        control
1354     C_error (np.ndarray): Covariance matrix for the control error
1355     x0 (np.ndarray): Initial ocean state
1356     timesteps (int): Number of timesteps for simulation
1357     ocean_states_observed (list): List of observed ocean states
1358     num_iters (int): Number of iterations for gradient descent
1359     step_size (float): Step size for gradient descent
1360     disp (bool, optional): If True, display progress. Defaults to False.
1361
1362     Returns:
1363     tuple: (f_adjust, losses, debug_vars)

```

```

1364         f_adjust (np.ndarray): Final adjustment to the atmospheric
1365             forcing field
1366         losses (dict): Dictionary of loss values over iterations
1367         debug_vars (dict): Dictionary of debug variables over iterations
1368     """
1369     return gradient_descent_template(M, F, f_true, f_guess, C_known,
1370                                     C_error,
1371                                     x0, timesteps,
1372                                     ocean_states_observed, num_iters,
1373                                     step_size,
1374                                     simple_gradient_update_rule, [],
1375                                     disp)
1376
1377 def cholesky_update_rule(curr_iter, s, step_size, f_adjust, C_error):
1378     """
1379     Computes the update step using Cholesky decomposition of the error
1380     covariance matrix.
1381
1382     Args:
1383     curr_iter (int): Current iteration number (unused in this function)
1384     s (np.ndarray): Gradient of the loss with respect to the forcing
1385                     field
1386     step_size (float): Step size for gradient descent
1387     f_adjust (np.ndarray): Current adjustment to the forcing field (
1388                           unused in this function)
1389     C_error (np.ndarray): Covariance matrix for the control error
1390
1391     Returns:
1392     np.ndarray: Update step for the forcing field adjustment
1393
1394     Notes:
1395     Applies the Cholesky decomposition L of C_error to s, then rescales
1396     the result
1397     to match the original gradient's magnitude.
1398     """
1399     # Apply cholesky decomposition L : C_error = L @ L.T
1400     L = np.linalg.cholesky(C_error)
1401     cholesky_s = L @ s
1402
1403     # Rescale so magnitude is the same
1404     rescaled_cholesky_s = cholesky_s * (np.linalg.norm(s) / np.linalg.
1405                                         norm(cholesky_s))
1406     step = - step_size * rescaled_cholesky_s
1407
1408     return step

```

```

1402
1403 def cholesky_gradient_descent(M, F, f_true, f_guess, C_known, C_error,
1404     # World parameters
1405
1406     x0, timesteps,
1407
1408     #
1409     Simulation parameters
1410     ocean_states_observed, num_iters, step_size
1411     , # Optimization parameters
1412     disp=False):
1413
1414     #
1415     Optimization method
1416
1417     """
1418     Implements gradient descent using Cholesky decomposition for
1419     optimizing the atmospheric forcing field.
1420
1421     Args:
1422     M (np.ndarray): Model matrix
1423     F (float): Forcing coefficient
1424     f_true (np.ndarray): True atmospheric forcing field
1425     f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1426     C_known (np.ndarray): Covariance matrix for the known portion of the
1427     control
1428     C_error (np.ndarray): Covariance matrix for the control error
1429     x0 (np.ndarray): Initial ocean state
1430     timesteps (int): Number of timesteps for simulation
1431     ocean_states_observed (list): List of observed ocean states
1432     num_iters (int): Number of iterations for gradient descent
1433     step_size (float): Step size for gradient descent
1434     disp (bool, optional): If True, display progress. Defaults to False.
1435
1436     Returns:
1437     tuple: (f_adjust, losses, debug_vars)
1438     f_adjust (np.ndarray): Final adjustment to the atmospheric
1439     forcing field
1440     losses (dict): Dictionary of loss values over iterations
1441     debug_vars (dict): Dictionary of debug variables over iterations
1442     """
1443     return gradient_descent_template(M, F, f_true, f_guess, C_known,
1444     C_error,
1445
1446     x0, timesteps,
1447     ocean_states_observed, num_iters,
1448     step_size,
1449     cholesky_update_rule, [C_error],
1450     disp)
1451
1452 def cov_constraint_J_update_rule(curr_iter, s, step_size, f_adjust,

```

```

C_error, weight_cov_term):
    """
    Computes the update step using a covariance constraint on the loss
        function.

    Args:
    curr_iter (int): Current iteration number (unused in this function)
    s (np.ndarray): Gradient of the loss with respect to the forcing
        field
    step_size (float): Step size for gradient descent
    f_adjust (np.ndarray): Current adjustment to the forcing field
    C_error (np.ndarray): Covariance matrix for the control error
    weight_cov_term (float): Weight for the covariance constraint term

    Returns:
    np.ndarray: Update step for the forcing field adjustment

    Notes:
    Adds a weighted covariance constraint term to the original gradient,
    then rescales the result to match the original gradient's magnitude.
    """
    cov_term_grad = 2 * np.linalg.inv(C_error) @ f_adjust # Covariance
        term

    s_prime = s + weight_cov_term * cov_term_grad # Gradient of J' with
        respect to the forcing field

    norm_s_prime = s_prime * (np.linalg.norm(s) / np.linalg.norm(s_prime)
        ) # Rescale magnitude to match original

    return -step_size * norm_s_prime

def cov_constraint_J_gradient_descent(M, F, f_true, f_guess, C_known,
    C_error, # World parameters
        x0, timesteps,

        # Simulation parameters
    ocean_states_observed, num_iters,
        step_size, # Optimization
        parameters
    weight_cov_term, disp=False):
        # Optimization
        method

    """
    Implements gradient descent with a covariance constraint for
        optimizing the atmospheric forcing field.

```



```

1470     Args:
1471     M (np.ndarray): Model matrix
1472     F (float): Forcing coefficient
1473     f_true (np.ndarray): True atmospheric forcing field
1474     f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1475     C_known (np.ndarray): Covariance matrix for the known portion of the
        control
1476     C_error (np.ndarray): Covariance matrix for the control error
1477     x0 (np.ndarray): Initial ocean state
1478     timesteps (int): Number of timesteps for simulation
1479     ocean_states_observed (list): List of observed ocean states
1480     num_iters (int): Number of iterations for gradient descent
1481     step_size (float): Step size for gradient descent
1482     weight_cov_term (float): Weight for the covariance constraint term
1483     disp (bool, optional): If True, display progress. Defaults to False.
1484
1485     Returns:
1486     tuple: (f_adjust, losses, debug_vars)
1487         f_adjust (np.ndarray): Final adjustment to the atmospheric
            forcing field
1488         losses (dict): Dictionary of loss values over iterations
1489         debug_vars (dict): Dictionary of debug variables over iterations
1490     """
1491     return gradient_descent_template(M, F, f_true, f_guess, C_known,
        C_error,
1492                                     x0, timesteps,
1493                                     ocean_states_observed, num_iters,
1494                                     step_size,
1495                                     cov_constraint_J_update_rule, [
        C_error, weight_cov_term], disp)
1496
1497 def dan_update_rule(curr_iter, s, step_size, f_adjust, C_error):
1498     """
1499     Computes the update step using Dan's method for improving the
        Mahalanobis distance.
1500
1501     Args:
1502     curr_iter (int): Current iteration number (unused in this function)
1503     s (np.ndarray): Gradient of the loss with respect to the forcing
        field
1504     step_size (float): Step size for gradient descent
1505     f_adjust (np.ndarray): Current adjustment to the forcing field (
        unused in this function)
1506     C_error (np.ndarray): Covariance matrix for the control error
1507
1508     Returns:
1509     np.ndarray: Update step for the forcing field adjustment

```

Notes:

Modifies the gradient direction to improve the Mahalanobis distance of u_i (update rule) while maintaining the desired improvement in the loss function J .

"""

$u_{i_simple_gd} = -step_size * s$ # Pre-dan step

$\Delta = s.T @ (u_{i_simple_gd})$ # Compute desired improvement of J

$new_vec = (C_error @ s) / (s.T @ C_error @ s)$ # Direction modified to improve Mahalanobis distance

return $\Delta * new_vec$

```
def dan_gradient_descent(M, F, f_true, f_guess, C_known, C_error, #
    World parameters
                        x0, timesteps, #
                        Simulation parameters
                        ocean_states_observed, num_iters, step_size, #
                        Optimization parameters
                        disp=False): #
    Optimization method

    """
    Implements gradient descent using Dan's method for optimizing the
    atmospheric forcing field.

    Args:
    M (np.ndarray): Model matrix
    F (float): Forcing coefficient
    f_true (np.ndarray): True atmospheric forcing field
    f_guess (np.ndarray): Initial guess for the atmospheric forcing field
    C_known (np.ndarray): Covariance matrix for the known portion of the
        control
    C_error (np.ndarray): Covariance matrix for the control error
    x0 (np.ndarray): Initial ocean state
    timesteps (int): Number of timesteps for simulation
    ocean_states_observed (list): List of observed ocean states
    num_iters (int): Number of iterations for gradient descent
    step_size (float): Step size for gradient descent
    disp (bool, optional): If True, display progress. Defaults to False.

    Returns:
    tuple: (f_adjust, losses, debug_vars)
        f_adjust (np.ndarray): Final adjustment to the atmospheric
            forcing field
        losses (dict): Dictionary of loss values over iterations
        debug_vars (dict): Dictionary of debug variables over iterations
```

```

1547 """
1548 return gradient_descent_template(M, F, f_true, f_guess, C_known,
1549                                 C_error,
1550                                 x0, timesteps,
1551                                 ocean_states_observed, num_iters,
1552                                 step_size,
1553                                 dan_update_rule, [C_error], disp)
1554
1555 def dan_modified_update_rule(curr_iter, s, step_size, f_adjust, C_error):
1556     """
1557     Computes the update step using a modified version of Dan's method for
1558     improving the Mahalanobis distance.
1559
1560     Args:
1561     curr_iter (int): Current iteration number (unused in this function)
1562     s (np.ndarray): Gradient of the loss with respect to the forcing
1563                     field
1564     step_size (float): Step size for gradient descent
1565     f_adjust (np.ndarray): Current adjustment to the forcing field
1566     C_error (np.ndarray): Covariance matrix for the control error
1567
1568     Returns:
1569     np.ndarray: Update step for the forcing field adjustment
1570
1571     Notes:
1572     Modifies the gradient direction to improve the Mahalanobis distance
1573     while
1574     maintaining the desired improvement in the loss function J. This
1575     version
1576     intends to improve Mahalanobis distance of  $a_i + u_i$  ( $f\_adjust + \text{gradient}$ 
1577     update), instead of just  $u_i$ .
1578     """
1579     ui_simple_gd = -step_size * s # Pre-dan step
1580     delta = s.T @ (ui_simple_gd) # Compute desired improvement of J
1581
1582     new_vec = (C_error @ s) / (s.T @ C_error @ s) # Direction modified
1583             to improve Mahalanobis distance
1584     vec_scale = delta + s.T @ f_adjust
1585
1586     return -f_adjust + vec_scale * new_vec
1587
1588 def dan_modified_gradient_descent(M, F, f_true, f_guess, C_known, C_error
1589 ,           # World parameters
1590
1591             x0, timesteps,
1592
1593             #
1594             Simulation parameters

```

```

1583         ocean_states_observed, num_iters,
1584             step_size, # Optimization
1585             parameters
1586         disp=False):
1587
1588             #
1589             Optimization method
1590
1591     """
1592     Implements gradient descent using a modified version of Dan's method
1593     for optimizing the atmospheric forcing field.
1594
1595     Args:
1596     M (np.ndarray): Model matrix
1597     F (float): Forcing coefficient
1598     f_true (np.ndarray): True atmospheric forcing field
1599     f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1600     C_known (np.ndarray): Covariance matrix for the known portion of the
1601         control
1602     C_error (np.ndarray): Covariance matrix for the control error
1603     x0 (np.ndarray): Initial ocean state
1604     timesteps (int): Number of timesteps for simulation
1605     ocean_states_observed (list): List of observed ocean states
1606     num_iters (int): Number of iterations for gradient descent
1607     step_size (float): Step size for gradient descent
1608     disp (bool, optional): If True, display progress. Defaults to False.
1609
1610     Returns:
1611     tuple: (f_adjust, losses, debug_vars)
1612         f_adjust (np.ndarray): Final adjustment to the atmospheric
1613             forcing field
1614         losses (dict): Dictionary of loss values over iterations
1615         debug_vars (dict): Dictionary of debug variables over iterations
1616     """
1617     return gradient_descent_template(M, F, f_true, f_guess, C_known,
1618         C_error,
1619         x0, timesteps,
1620         ocean_states_observed, num_iters,
1621         step_size,
1622         dan_modified_update_rule, [C_error],
1623         disp)
1624
1625 ### Gradient Descent Testing ###
1626
1627 def compare_gd_methods_once(M, F, f_true, f_guess, C_known, C_error,
1628     x0, timesteps,
1629     ocean_states_observed, num_iters, step_size,
1630     methods, disp=False):

```

```

1620 """
1621 Runs multiple methods of gradient descent on the same dataset and
1622     compares their performance.
1623
1624 Args:
1625 M (np.ndarray): Model matrix
1626 F (float): Forcing coefficient
1627 f_true (np.ndarray): True atmospheric forcing field
1628 f_guess (np.ndarray): Initial guess for the atmospheric forcing field
1629 C_known (np.ndarray): Covariance matrix for the known portion of the
1630     control
1631 C_error (np.ndarray): Covariance matrix for the control error
1632 x0 (np.ndarray): Initial ocean state
1633 timesteps (int): Number of timesteps for simulation
1634 ocean_states_observed (list): List of observed ocean states
1635 num_iters (int): Number of iterations for gradient descent
1636 step_size (float): Step size for gradient descent
1637 methods (list): List of gradient descent methods to compare. Each
1638     method is a list of the form
1639         ["Method Name", method_func, extra_params]
1640 disp (bool, optional): If True, display progress. Defaults to False.
1641
1642 Returns:
1643 dict: A dictionary where keys are method names and values are tuples
1644     containing:
1645     - f_adjust (np.ndarray): Final adjustment to the atmospheric
1646         forcing field
1647     - losses (dict): Dictionary of loss values over iterations
1648     - debug_vars (dict): Dictionary of debug variables over
1649         iterations
1650
1651 Notes:
1652 This function applies each specified gradient descent method to the
1653     same initial conditions
1654     and dataset, allowing for direct comparison of their performance.
1655 """
1656 results = {}
1657
1658 for method_name, method_func, extra_params in methods:
1659     if disp:
1660         print(f"Running method {method_name}")
1661
1662     f_adjust, losses, debug_vars = gradient_descent_template(M, F,
1663         f_true, f_guess, C_known, C_error,
1664         x0, timesteps,
1665         ocean_states_observed,
1666         num_iters,

```

```

1658                                     step_size,
                                     method_func,
                                     extra_params,
                                     disp)
1659     results[method_name] = (f_adjust, losses, debug_vars)
1660
1661     return results
1662
1663 def compare_gd_methods_many_times(nr, nc, dt, F, gamma, sigma,
1664     num_obs_per_timestep,
1665     num_timesteps, num_iters, step_size,
1666     C_known, C_error, methods, num_runs,
1667     disp=False):
1668     """
1669     Create many different sets of data.
1670     For each one, we will run each of our gradient descent methods.
1671     Once we finish, we average losses across all runs.
1672
1673     Args:
1674     nr (int): Number of rows in the grid
1675     nc (int): Number of columns in the grid
1676     dt (float): Time step
1677     F (float): Forcing parameter
1678     gamma (float): Proportion of the control vector that is correct
1679     sigma (float): Standard deviation of observation noise
1680     num_obs_per_timestep (int): Number of observations per timestep
1681     num_timesteps (int): Number of timesteps
1682     num_iters (int): Number of iterations of gradient descent
1683     step_size (float): Step size for gradient descent
1684     methods (list): List of gradient descent methods to compare
1685     num_runs (int): Number of times to run the whole optimization process
1686     disp (bool): If True, print progress
1687
1688     Returns:
1689     dict: Dictionary containing averaged losses and debug variables for
1690     each method
1691     """
1692     # Initialize results dictionary
1693     losses = {method[0]: copy.deepcopy(losses_template)
1694         for method in methods}
1695     debug_vars = {method[0]: copy.deepcopy(possible_debug_vars)
1696         for method in methods}
1697
1698     for run in range(num_runs):
1699         if disp:
1700             print(f"Run {run + 1}/{num_runs}")

```

```

1699
1700     # Generate world
1701     C_control, x0, _, M = generate_world(nr, nc, dt, F)
1702     C_known, C_error = C_control * gamma, C_control * (1-gamma)
1703
1704     f_true, f_guess = generate_true_and_first_guess_field(C_known,
1705         C_error, nr, nc)
1706
1707     # Run the simulation with the true and guessed control vector
1708     _, real_state_over_time = compute_affine_time_evolution_simple(
1709         x0, M, F*f_true, num_timesteps)
1710
1711     observed_state_over_time_2d = observe_over_time(
1712         real_state_over_time, sigma,
1713         num_obs_per_timestep, nr, nc)
1714
1715     observed_state_over_time = [np.reshape(observed_state_2d, (nr*nc,
1716         1))
1717         for observed_state_2d in
1718             observed_state_over_time_2d]
1719
1720     # Run each method
1721     for method_name, method_func, extra_params in methods:
1722         if disp:
1723             print(f" Method: {method_name}")
1724
1725         results = compare_gd_methods_once(M, F, f_true, f_guess,
1726             C_known, C_error,
1727             x0, num_timesteps,
1728             observed_state_over_time,
1729             num_iters, step_size,
1730             [[method_name,
1731                 method_func,
1732                 extra_params]], disp)
1733
1734     # Include new losses
1735     for method_name, (_, method_losses, method_debug_vars) in
1736         results.items():
1737         for loss_name, loss_list in losses[method_name].items():
1738             loss_list.append(method_losses[loss_name])
1739
1740         for debug_name, debug_list in debug_vars[method_name].
1741             items():
1742             debug_list.append(method_debug_vars[debug_name])

```

```
1732
1733     # Average losses
1734
1735     for method_name, method_losses in losses.items():
1736         for loss_name, loss_list in method_losses.items():
1737             losses[method_name][loss_name] = np.mean(loss_list, axis=0)
1738
1739     for method_name, method_debug_vars in debug_vars.items():
1740         for debug_name, debug_list in method_debug_vars.items():
1741             debug_vars[method_name][debug_name] = np.mean(debug_list,
1742                                                             axis=0)
1743
1744     return losses, debug_vars
1745
1746 ### Gradient Descent Visualization ###
1747
1748 def plot_losses(losses_many, num_obs_per_timestep, step_size,
1749               num_timesteps, num_iters, min_iter=None, max_iter=None):
1750     """
1751     Plots the losses for multiple gradient descent methods over
1752     iterations.
1753
1754     Args:
1755     losses_many (dict): Dictionary of losses for each method.
1756                         Keys are method names, values are dictionaries
1757                         containing losses.
1758     num_obs_per_timestep (int): Number of observations per timestep
1759     step_size (float): Step size used in gradient descent
1760     num_timesteps (int): Number of timesteps in the simulation
1761     num_iters (int): Number of iterations of gradient descent
1762     min_iter (int): Lowest plotted iter (default: None, plots from the
1763                     beginning)
1764     max_iter (int): Highest plotted iter (default: None, plots until the
1765                     end)
1766
1767     Returns:
1768     None: This function displays the plot using matplotlib.pyplot.show()
1769
1770     Notes:
1771     Creates a 2x2 grid of plots:
1772     1. Ocean misfit
1773     2. Atmosphere loss
1774     3. Control adjust Mahalanobis distance
1775     4. J' (combined loss for covariance constraint method, ocean loss for
1776         others)
1777
1778     Each plot shows the evolution of the respective loss over iterations
```



```

    for all methods.
1772 """
1773 fig, axs = plt.subplots(2, 2, figsize=(10, 10))
1774
1775 loss_funcs = ["$\sum_t (Ex(t)-y(t))^T (Ex(t)-y(t))$",
1776               "$\sum_t (f_i(t)-f_{\text{true}}(t))^T (f_i(t)-f_{\text{true}}(t))$",
1777               "$a_i^T C^{-1} a_i$"]
1778
1779 # Determine the range of iterations to plot
1780 min_iter = 0 if min_iter is None else max(0, min_iter)
1781 max_iter = num_iters if max_iter is None else min(num_iters, max_iter)
1782
1783 plot_range = slice(min_iter, max_iter)
1784
1785 for i, (loss_name, ax, func) in enumerate(zip(["ocean_misfit", "
1786         atmosphere_misfit", "mahalanobis(covariance similarity)"], axs.
1787         flatten(), loss_funcs)):
1788     for method_name, losses_dict in losses_many.items():
1789         ax.plot(range(min_iter, max_iter), losses_dict[loss_name][
1790             plot_range], label=method_name)
1791         ax.set_xlabel("Iteration $i$")
1792         ax.set_ylabel(loss_name+" loss: "+func)
1793         ax.legend()
1794         ax.set_title(f"{loss_name}: "+func)
1795
1796     # Set integer ticks on x-axis
1797     ax.xaxis.set_major_locator(plt.MaxNLocator(integer=True))
1798
1799 # Fourth plot: ocean_misfit + mahalanobis if using covariance control
1800 adjust, just ocean otherwise
1801 ax = axs[1, 1]
1802 for method_name, losses_dict in losses_many.items():
1803     if method_name == r"Covariance Constraint J Gradient Descent":
1804         combined_loss = [o + c for o, c in zip(losses_dict["
1805             ocean_misfit"], losses_dict["mahalanobis(covariance
1806             similarity)"])]
1807         ax.plot(range(min_iter, max_iter), combined_loss[plot_range],
1808             label=method_name)
1809     else:
1810         ax.plot(range(min_iter, max_iter), losses_dict["ocean_misfit"
1811             ][plot_range], label=method_name)
1812
1813 ax.set_xlabel("Iteration $i$")
1814 ax.set_ylabel("$J'$")
1815 ax.legend()
1816 ax.set_title("$J'$")

```

```

1808
1809 # Set integer ticks on x-axis for the fourth plot
1810 ax.xaxis.set_major_locator(plt.MaxNLocator(integer=True))
1811
1812 fig.suptitle(f"Gradient Descent Variants: num_obs={
1813     num_obs_per_timestep}, step_size={step_size}, num_timesteps={
1814     num_timesteps}, num_iters={num_iters}")
1815
1816 plt.tight_layout()
1817 plt.show()
1818
1819 def plot_debug(debug_vars, min_iter=None, max_iter=None, tickwidth=1,
1820     vlines = []):
1821     """
1822     Plot 9 chosen debug variables in a 3x3 grid for each method.
1823
1824     Args:
1825     debug_vars (dict): Dictionary containing debug variables for each
1826         method
1827     min_iter (int, optional): Start index for plotting. If None, starts
1828         from the beginning.
1829     max_iter (int, optional): End index for plotting. If None, plots
1830         until the end.
1831     tickwidth (int): Width between ticks on x-axis
1832     vlines (list): List of vertical lines to add to the plot
1833
1834     Returns:
1835     None: Displays the plot
1836     """
1837     # Determine the actual range of iterations
1838     all_iters = next(iter(debug_vars.values()))["Norm of s_i"]
1839     total_iters = len(all_iters)
1840
1841     # Set min_iter and max_iter if they are None
1842     min_iter = 0 if min_iter is None else max(0, min_iter)
1843     max_iter = total_iters if max_iter is None else min(total_iters,
1844         max_iter)
1845
1846     fig, axs = plt.subplots(3, 3, figsize=(15, 15))
1847     fig.suptitle(f"Debug Variables: Iterations {min_iter} to {max_iter}")
1848
1849     # List of debug variables to plot
1850     plot_vars = [
1851         "Norm of s_i",
1852         "Expected Delta J w simple gd",
1853         "Expected Delta J w update rule",
1854         "Actual Delta J",

```

```
1848     "Norm of simple gd ui",
1849     "Norm of update rule ui",
1850     "Normalized dot product $a_i$ and $u_i$",
1851     "Normalized $-Cs_i \cdot a_i$",
1852     "Norm of $a_i$",
1853 ]
1854
1855 for i, (var_name, ax) in enumerate(zip(plot_vars, axs.flatten())):
1856     for method_name, method_debug_vars in debug_vars.items():
1857         if var_name in method_debug_vars:
1858             plot_data = method_debug_vars[var_name][min_iter:max_iter
1859             ]
1860             ax.plot(range(min_iter, max_iter), plot_data, "-.", label
1861                     =method_name)
1862
1863     ax.set_xlabel("Iteration $i$")
1864     ax.set_ylabel(var_name)
1865     ax.legend()
1866     ax.set_title(var_name)
1867     ax.grid(True)
1868
1869     # Set x-axis ticks to reflect the actual iteration numbers
1870     ticks = range(min_iter, max_iter, tickwidth)
1871     ax.set_xticks(ticks)
1872     ax.set_xticklabels(ticks)
1873
1874     # Add vertical lines if specified
1875     for vline in vl_lines:
1876         if min_iter <= vline < max_iter:
1877             ax.axvline(x=vline, color='r', linestyle='--', alpha=0.5)
1878
1879 plt.tight_layout()
1880 plt.show()
```